


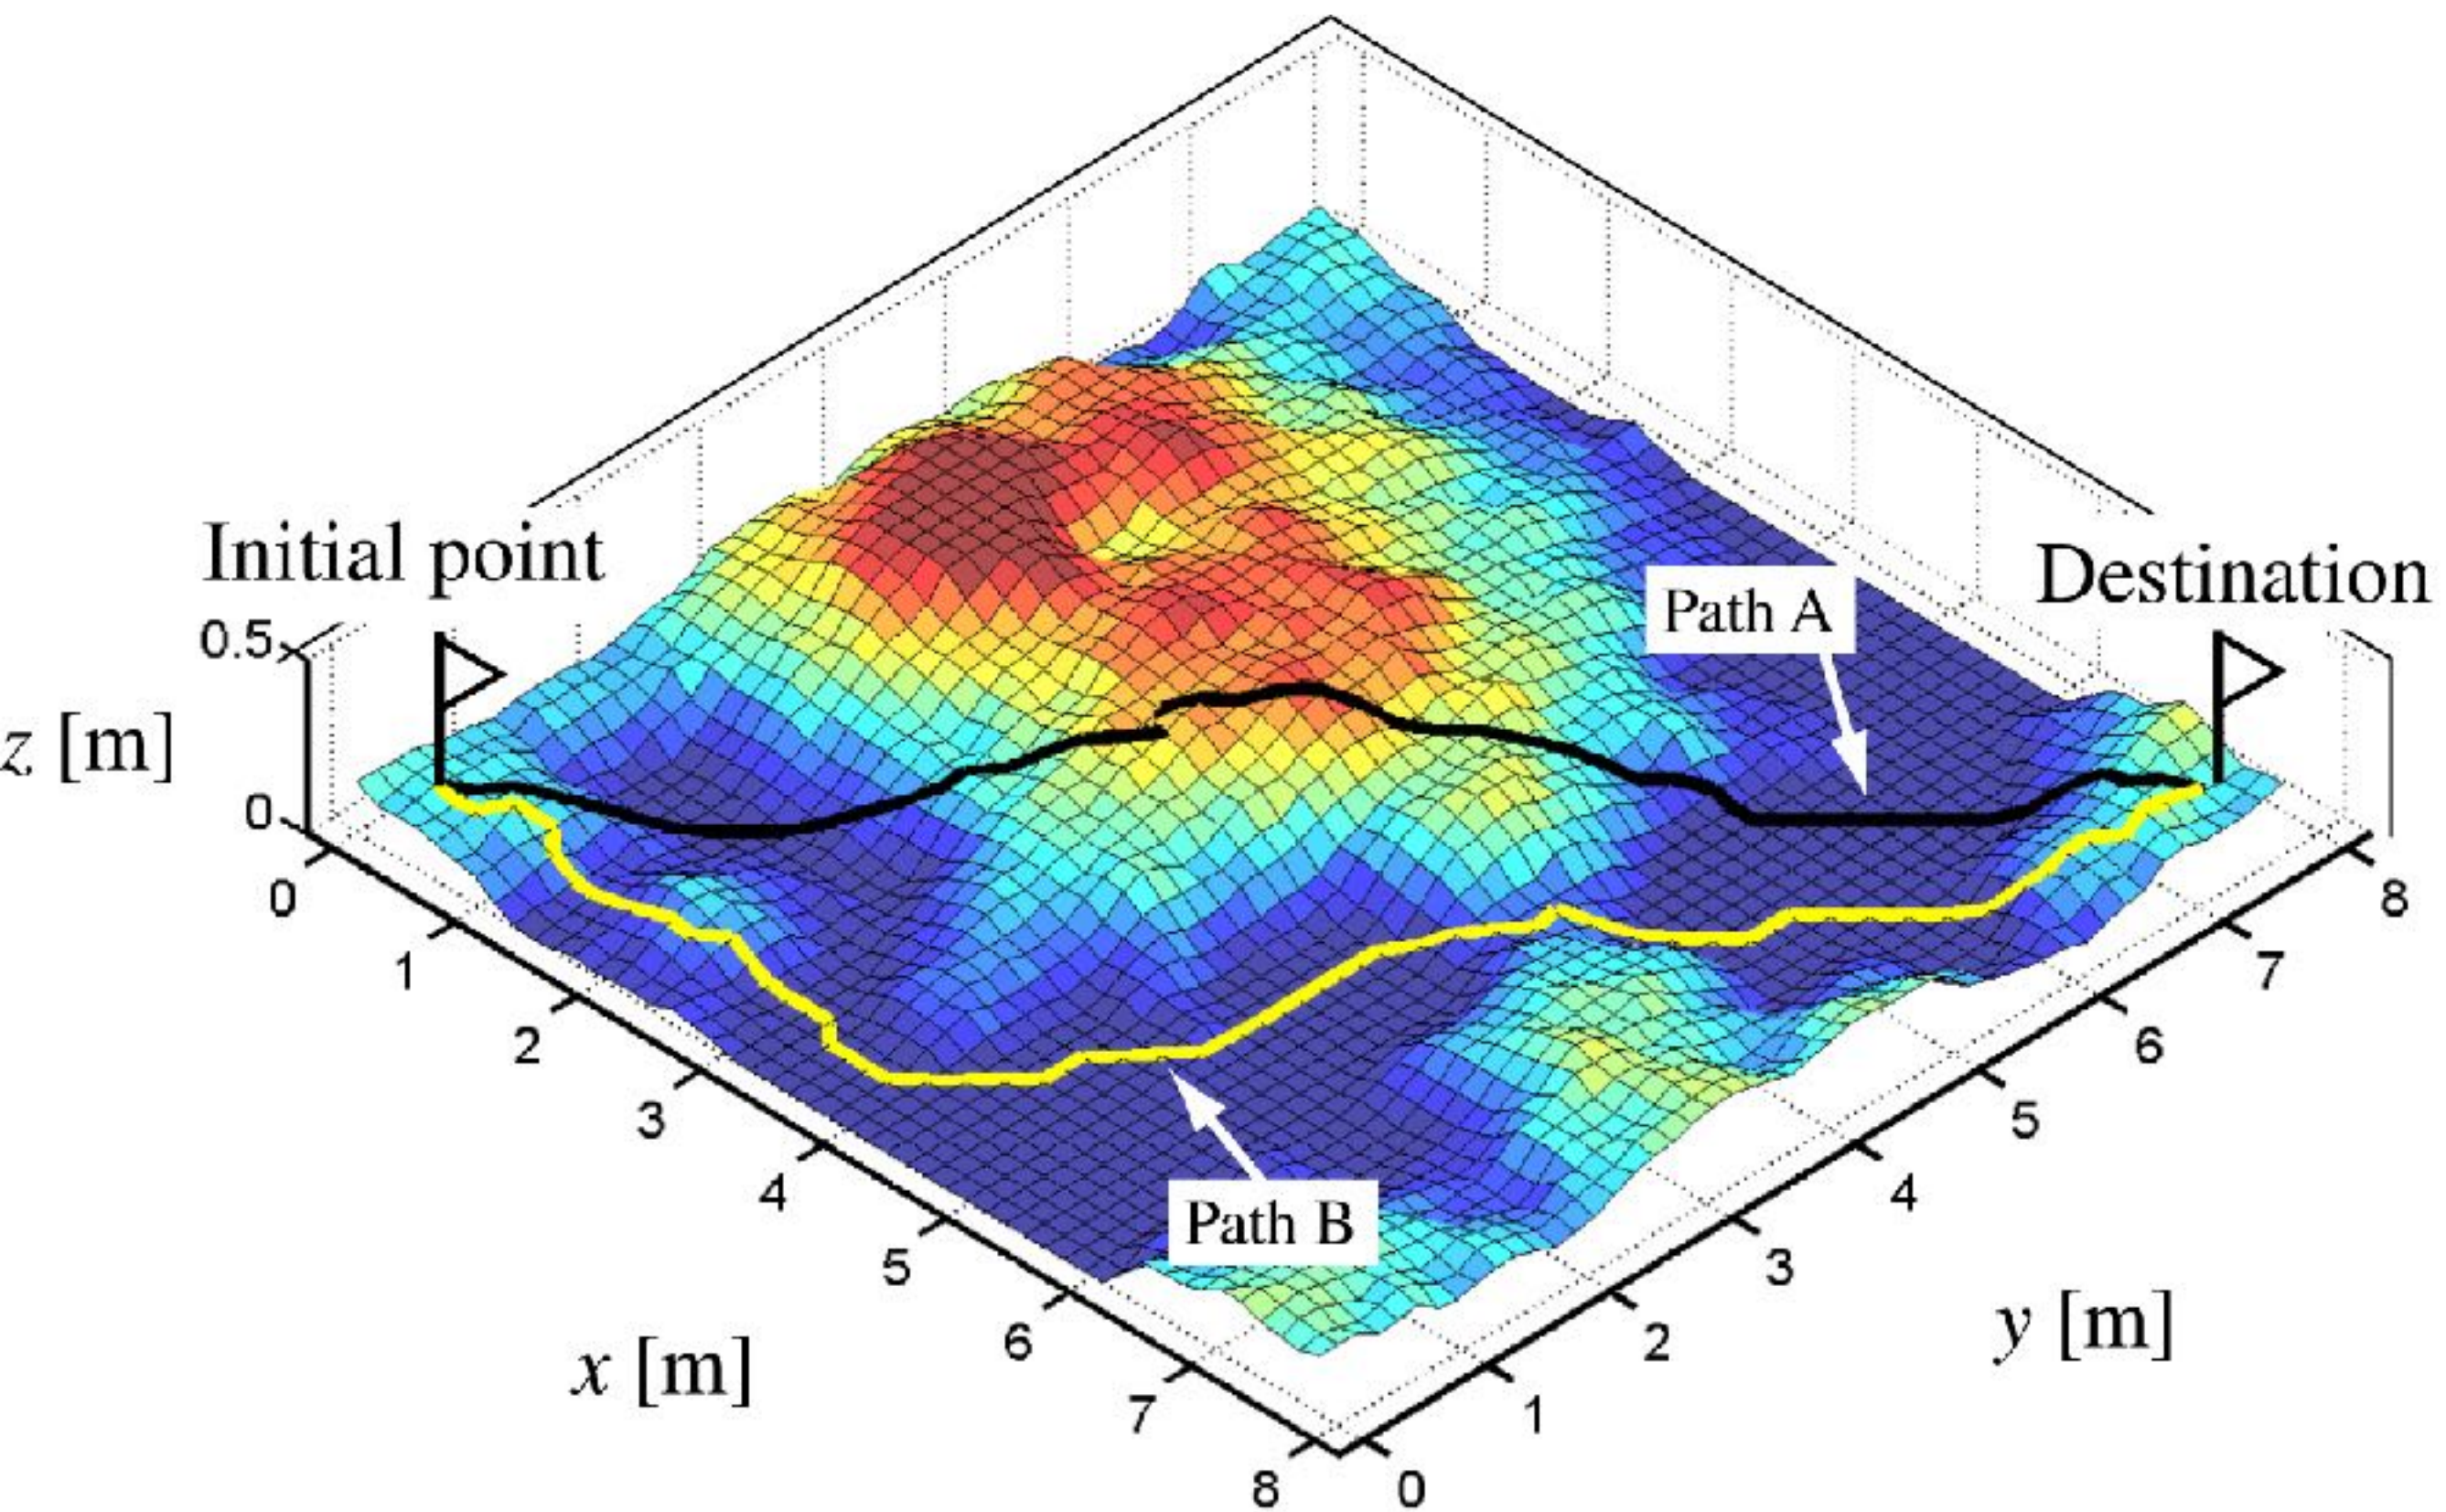
ARTIFICIAL INTELLIGENCE

Problem Solving

PATH PLANNING



Retrieved from: https://www.mathworks.com/help/examples/robotics_product/win64/PathPlanningExample_03.png



Retrieved from: <http://www.astro.mech.tohoku.ac.jp/~ishigami/research/image/path.jpg>

STATE SPACE

BEFORE SOLVING THE PROBLEM

You are given a problem to solve.

What do you need to do first?

- For example, when you want to complete a jigsaw puzzle, what do you do first?

BEFORE SOLVING THE PROBLEM

What do you need to do first?



Retrieved from https://lh6.ggpht.com/v_0cglTjMUfOzijN1Pq5plteMBjmo7NXhUmx8kzkGD1irQ-dH2-_s-Y7hhes3q-KjxJz=h900

BREAKING DOWN THE PROBLEM

Identify:

- Initial State
 - What does the problem look like in the beginning?
- Current State
 - What does the problem look like right now?
- Goal State
 - What is the final goal?
- Actions we can take
 - What can we act on the problem?
- Result after the action is taken
 - What happens after we perform the action.

STATE SPACE

$S = \langle S, A, \text{Action}(s), \text{Result}(s,a), \text{Cost}(s,a) \rangle$

S = Set of all states

A = Set of all action


$\text{Action}(s)$ = All valid actions at state s

$\text{Result}(s,a)$ = State after taking action a at state s

$\text{Cost}(s,a)$ = Cost associated when taking action a at state s

PROBLEM FORMATION – ROBOT NAVIGATION

- Given that a robot starts at the green square and tries to go to the red square. The robot can move up, down, left, and right. Black squares represents impassable terrain. White squares represent passable terrain. And blue squares represent hard-to-navigate terrain. How do we set up this state space?



SEARCHING FOR SOLUTION



SEARCHING

We have

- State Spaces
- Initial State
- Goal State

Our job is to find a way to move from Initial State to Goal State


We also need to find a way to represent these actions

- Tree
- Graph
- Grid

ADVANCED DATA STRUCTURE

TREE

- Nodes with multiple links to other nodes.



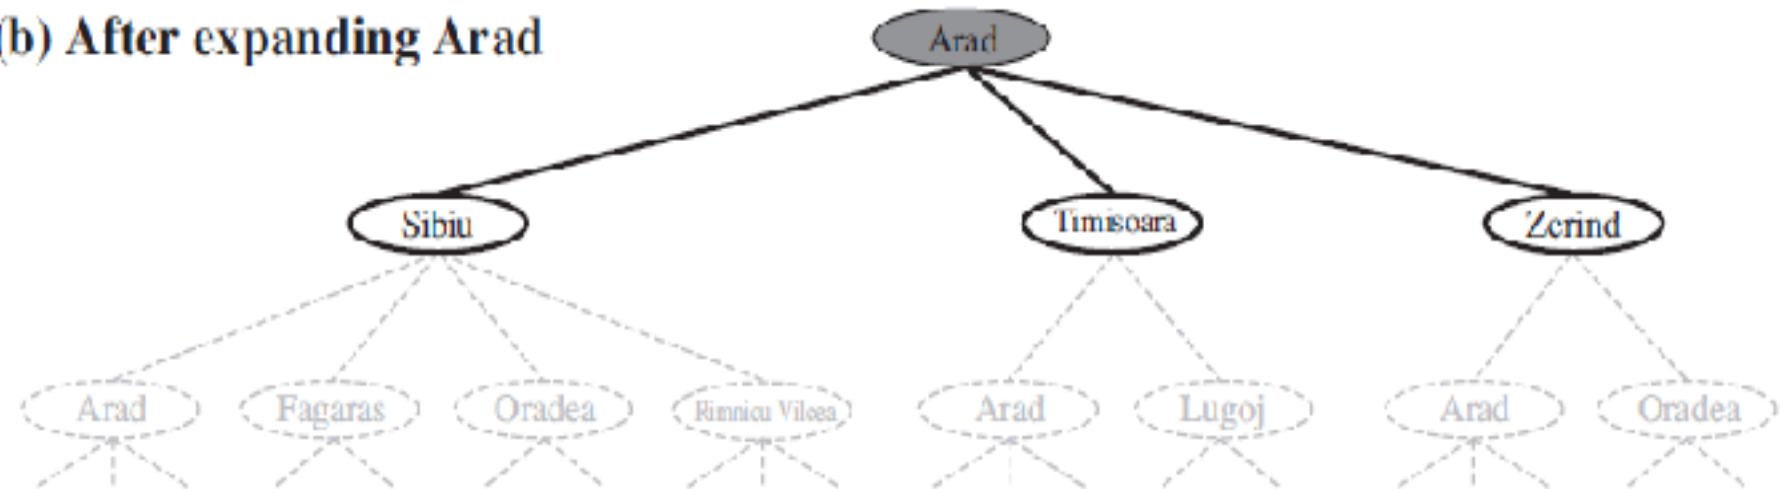
TREE

.....


(a) The initial state



(b) After expanding Arad




(c) After expanding Sibiu




GRAPH

.....



GRAPH

- A graph $G = (V, E)$
 - V = set of vertices
 - E = set of edges



➤

SEARCHING ALGORITHM

INFORMED SEARCH VS UNINFORMED SEARCH

Uninformed Search


- Key Idea:
 - The algorithm doesn't know if state s is close to the goal state or not.

Informed Search

- Key Idea:
 - The algorithm know if state s is better than other state in order to reach the goal.


UNINFORMED SEARCH

- A robot at both states A and B can perform 2 actions: West, and South.
- It doesn't know if B is closer to the goal state or not.



INFORMED SEARCH

- The robot knows that state B is closer to the goal state than state A.
- It can take this information in to efficiently solve the problem.



UNINFORMED SEARCH

UNINFORMED SEARCH

- Key Idea:
 - The algorithm doesn't know if state s is close to the goal state or not.
- Definition
 - The strategies have no additional information about states beyond that provided in the problem definition.

BREADTH-FIRST SEARCH

Uninformed Search

BREADTH-FIRST SEARCH

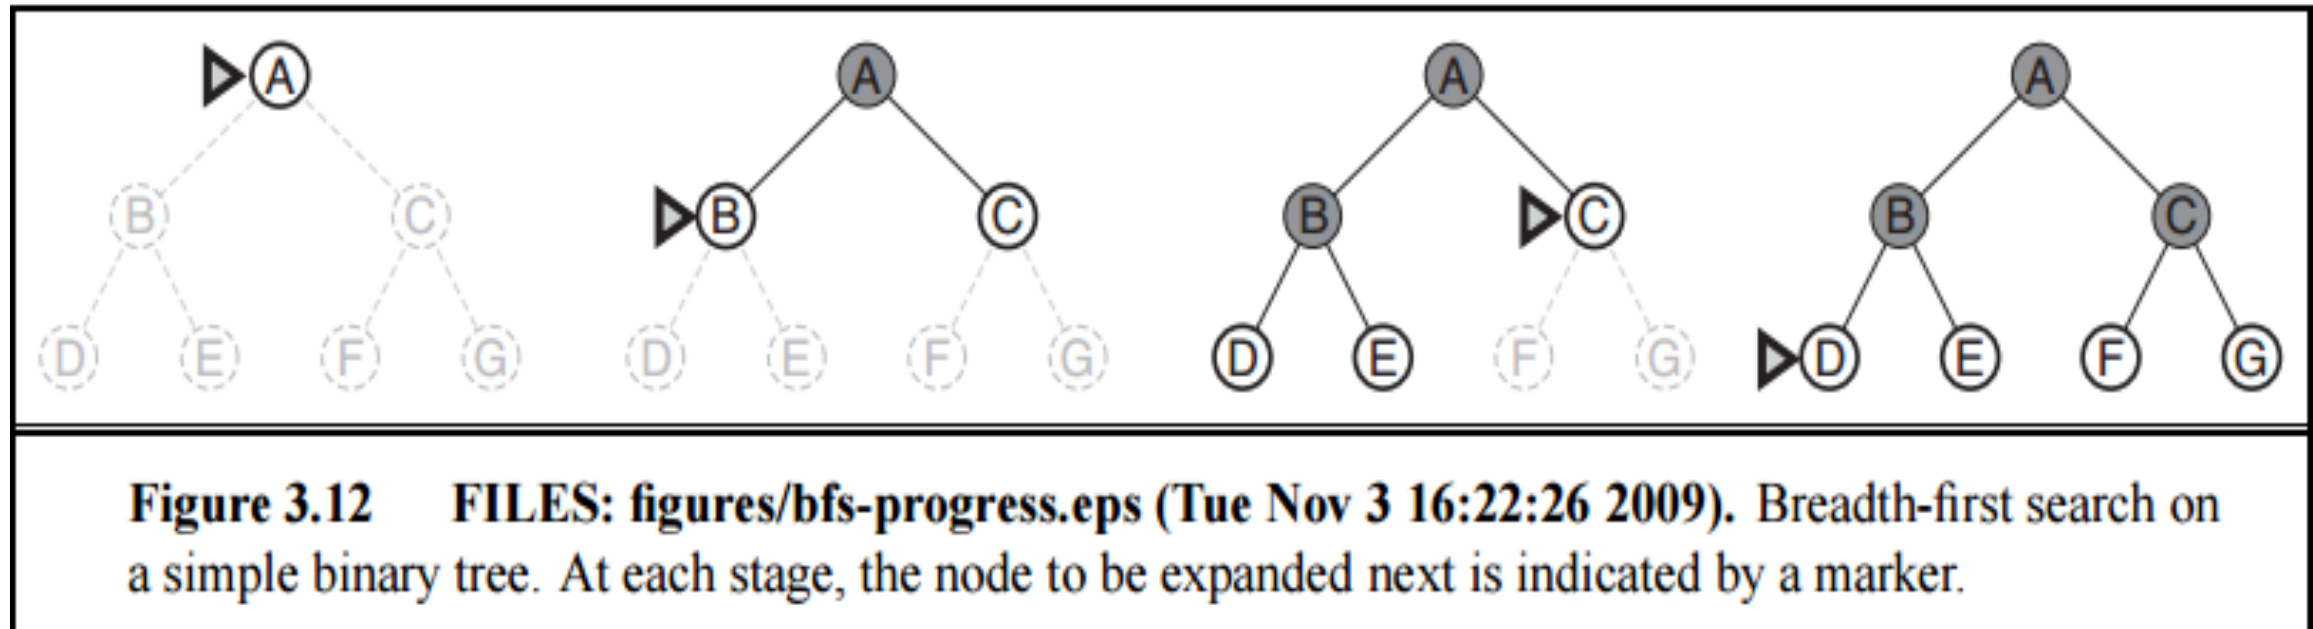
Key Idea

- Expand first, then search, then expand, ...

Definition

- A simple strategy in which the root node is expanded first, then all the successors of the root node are expanded next, then their successors, and so on.

BFS EXAMPLE



BFS: THE ALGORITHM

.....

```
function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier  $\leftarrow$  a FIFO queue with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the shallowest node in frontier */
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier  $\leftarrow$  INSERT(child, frontier)
```

Figure 3.11 Breadth-first search on a graph.

PROBLEMS WITH BREADTH-FIRST SEARCH

Given branching factor = 10, processing power = 1M nodes/sec, 1000 bytes/node:

Depth	Nodes	Time	Memory
2	110	.11 ms	107 kB
4	11110	11 ms	10.6 MB
10	10^{10}	3 hours	10 TB
16	10^{16}	350 years	10 EB

Required time and space increases exponentially.

BFS QUIZ

	1	2	3	4	5
1	s				
2					
3					
4					
5		g			

Starting state = (1,1)

Goal state = (2,5)

a. จงวาด tree ที่เกิดจากการทำ path planning โดยใช้ Breadth-first Search ให้ใส่ตัวเลขที่แต่ละกิ่ง (branch) เพื่อแสดงลำดับการแตกกิ่ง

b. Path ของการทำ path planning จาก Breath-first Search คืออะไร

UNIFORM-COST SEARCH

Uninformed Search

UNIFORM-COST SEARCH


Key Idea:

- Don't care about numbers of paths. Care about total cost.

Definition:

- Uniform-cost search expands the node n with the lowest path cost $g(n)$ first.

UCS EXAMPLE




UCS EXAMPLE

Initial State: Sibiu

Goal State: Bucharest

Start at Sibiu; cost = 0




UCS EXAMPLE

First, expand root node.

Sibui => Fagaras cost = 99

 => Rimnicu Vilcea cost = 80




UCS EXAMPLE

Next, expand the lower cost path.

Sibui => Fagara cost = 99

=> Rimnicu Vilcea=> Pitesti cost = 80 + 97 = 177




UCS EXAMPLE

Keep expanding the lower cost path.

Sibui \Rightarrow Fagaras \Rightarrow Bucharest cost = $99 + 211 = 310$

\Rightarrow Rimnicu Vilcea \Rightarrow Pitesti cost = 177




UCS EXAMPLE

Keep expanding the lower cost path.

Sibui => Fagaras => Bucharest cost = 310

 => Rimnicu Vilcea => Pitesti => Bucharest


 cost = 177 + 101 = 278



UCS EXAMPLE

Goal State is reached. Answer is the lowest cost path:

Sibui => Rimnicu Vilcea => Pitesti => Bucharest
cost = 278



UCS: THE ALGORITHM

function UNIFORM-COST-SEARCH(*problem*) **returns** a solution, or failure

node \leftarrow a node with STATE = *problem*.INITIAL-STATE, PATH-COST = 0

frontier \leftarrow a priority queue ordered by PATH-COST, with *node* as the only element

explored \leftarrow an empty set

loop do

if EMPTY?(*frontier*) **then return** failure

node \leftarrow POP(*frontier*) /* chooses the lowest-cost node in *frontier* */

if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)

 add *node*.STATE to *explored*

for each *action* **in** *problem*.ACTIONS(*node*.STATE) **do**

child \leftarrow CHILD-NODE(*problem*, *node*, *action*)

if *child*.STATE is not in *explored* or *frontier* **then**

frontier \leftarrow INSERT(*child*, *frontier*)

else if *child*.STATE is in *frontier* with higher PATH-COST **then**

 replace that *frontier* node with *child*

DEPTH-FIRST SEARCH

Uninformed Search




DEPTH-FIRST SEARCH

Key Idea:










- Expand the lowest node first..

Definition:

- Depth-first Search expands the deepest node in the current frontier of the search tree.



.....



PROBLEM WITH DEPTH-FIRST SEARCH

Searching in infinite spaces

- What if a child node of A is B and C, and a child node of B is A and D.
- We will have a tree that looks like:

DFS QUIZ

	1	2	3	4	5
1	s				
2					
3					
4					
5		g			

Starting state = (1,1)

Goal state = (2,5)

c. กำหนดให้ลำดับการแตก neighbor เป็นดังนี้: east, south, west, north, จงวาด tree ที่เกิดจากการทำ path planning โดยใช้ Depth-first Search ให้ใส่ตัวเลขที่แต่ละกิ่ง(branch) เพื่อแสดงลำดับการแตกกิ่ง

d. Path ของการทำ path planning จาก Depth-first Search คืออะไร

DEPTH-LIMITED SEARCH

Uninformed Search

DEPTH-LIMITED SEARCH

Key Idea:

- Same as Depth-first search, with limit on the depth (maximum level to expand)

Definition:

- Depth-limited Search is supplied with a predetermined depth limit l . Nodes at depth l are treated as if they have no successors.

DLS: THE ALGORITHM

.....

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff  
  return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)
```

```
function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff  
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)  
  else if limit = 0 then return cutoff  
  else  
    cutoff_occurred?  $\leftarrow$  false  
    for each action in problem.ACTIONS(node.STATE) do  
      child  $\leftarrow$  CHILD-NODE(problem, node, action)  
      result  $\leftarrow$  RECURSIVE-DLS(child, problem, limit - 1)  
      if result = cutoff then cutoff_occurred?  $\leftarrow$  true  
      else if result  $\neq$  failure then return result  
  if cutoff_occurred? then return cutoff else return failure
```

PROBLEMS WITH DLS

How do we pick the depth limit?

- We need to supply a depth limit l to the algorithm, so the algorithm terminates at some points.
- What if the goal state is at depth $l + 1$?

ITERATIVE DEEPENING DEPTH-FIRST SEARCH

Uninformed Search

ITERATIVE DEEPENING DEPTH-FIRST SEARCH

Key Idea:

- Iterating DFS with increasing depth limit.

Definition:

- Iterative Deepening Depth-first Search gradually increases the limit - first 0, then 1, then 2, and so on - until the goal is found.


IDDFS: EXAMPLE

.....







Limit = 0



Limit = 1




Limit = 2



IDDFS: EXAMPLE

Limit = 3




IDDFS: THE ALGORITHM

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result
```

OTHER ALGORITHM

Bidirectional Search

- Two simultaneous searches: one from start state, one from goal state
- Check if expanded nodes are overlapped



LAB TIME

- Problem Set 1 will consist of 3 Problem Solving minilab.
- Go to github.com/BawornsakS/FRA311
 - Access Problem Solving Lab 1 and follow the instruction.
 - You can write your answer and submit it via Google Classroom.

PATH PLANNING & OPTIMIZATION PROBLEM

OPTIMIZATION

- According to Google Translate, optimization is “the action of making the best or most effective use of a situation or resource.”
- Basically, finding the **best** moves to perform, so it can lead to the goal **efficiently**.

ADVANCED DATA STRUCTURE

PRIORITY QUEUE

- A queue (FIFO) with a priority.
 - We can construct it in a form of (element, priority).
- Given a queue:
 - (A,1), (B,4), (Z,7), (B, 10)
- If we want to add another node (C, 5)
 - Normal queue will output:
 - Priority queue will output:

PRIORITY QUEUE

- A queue (FIFO) with a priority.
 - We can construct it in a form of (element, priority).
- Given a queue:
 - (A,1), (B,4), (Z,7), (B, 10)
- If we want to add another node (C, 5)
 - Normal queue will output: (A,1), (B,4), (Z,7), (B, 10), (C, 5)
 - Priority queue will output:

PRIORITY QUEUE

- A queue (FIFO) with a priority.
 - We can construct it in a form of (element, priority).
- Given a queue:
 - (A,1), (B,4), (Z,7), (B, 10)
- If we want to add another node (C, 5)
 - Normal queue will output: (A,1), (B,4), (Z,7), (B, 10), (C, 5)
 - Priority queue will output: (A,1), (B,4), (C, 5), (Z,7), (B, 10)

HEURISTIC SEARCH

Informed Search

INFORMED SEARCH VS UNINFORMED SEARCH

Uninformed Search

- Key Idea:

The algorithm doesn't know if state s is close to the goal state or not.

Informed Search (Heuristic Search)

- Key Idea:

The algorithm know if state s is better than other state in order to reach the goal

HEURISTIC SEARCH

Key Idea:

The algorithm know if state s is better than other state in order to reach the goal

Definition:

Informed Search (Heuristic Search) uses problem-specific knowledge beyond the definition of the problem itself to determine the best path to the goal

BEYOND PROBLEMS SPECIFIC KNOWLEDGE

- Let's define what the problem specific knowledge are.
- In problem formation from the last lecture, what do we know?

BEYOND PROBLEMS SPECIFIC KNOWLEDGE

- Let's define what the problem specific knowledge are.
- In problem formation from the last lecture, what do we know?
 - Initial state
 - Goal state
 - All states
 - Actions
 - $\text{Result}(s,a)$
 - $\text{Cost}(s,a)$

BEYOND PROBLEMS SPECIFIC KNOWLEDGE


- In problem formation from the last lecture, what do we know?
 - Initial state, goal state, all states, actions, $\text{result}(s,a)$, $\text{cost}(s,a)$
- From these information, can we derive some additional information from them?

BEYOND PROBLEMS SPECIFIC KNOWLEDGE

- In problem formation from the last lecture, what do we know?
 - Initial state, goal state, all states, actions, $\text{result}(s,a)$, $\text{cost}(s,a)$
- From these information, can we derive some additional information from them?
- What if the world is a grid, and we want to find how far we are from the goal state?

BEYOND PROBLEMS SPECIFIC KNOWLEDGE

- If we want to know how far A and B are from the goal (red), how do we do it?




HEURISTIC FUNCTION

- For Informed Search, we use a Heuristic Function, $h(n)$, to determine the node to expand..
- Definition:
 - Heuristic function **estimates** cost of the cheapest path from the state at node n to a goal state.

BEYOND PROBLEMS SPECIFIC KNOWLEDGE

- What are the heuristic functions you can use on this map?

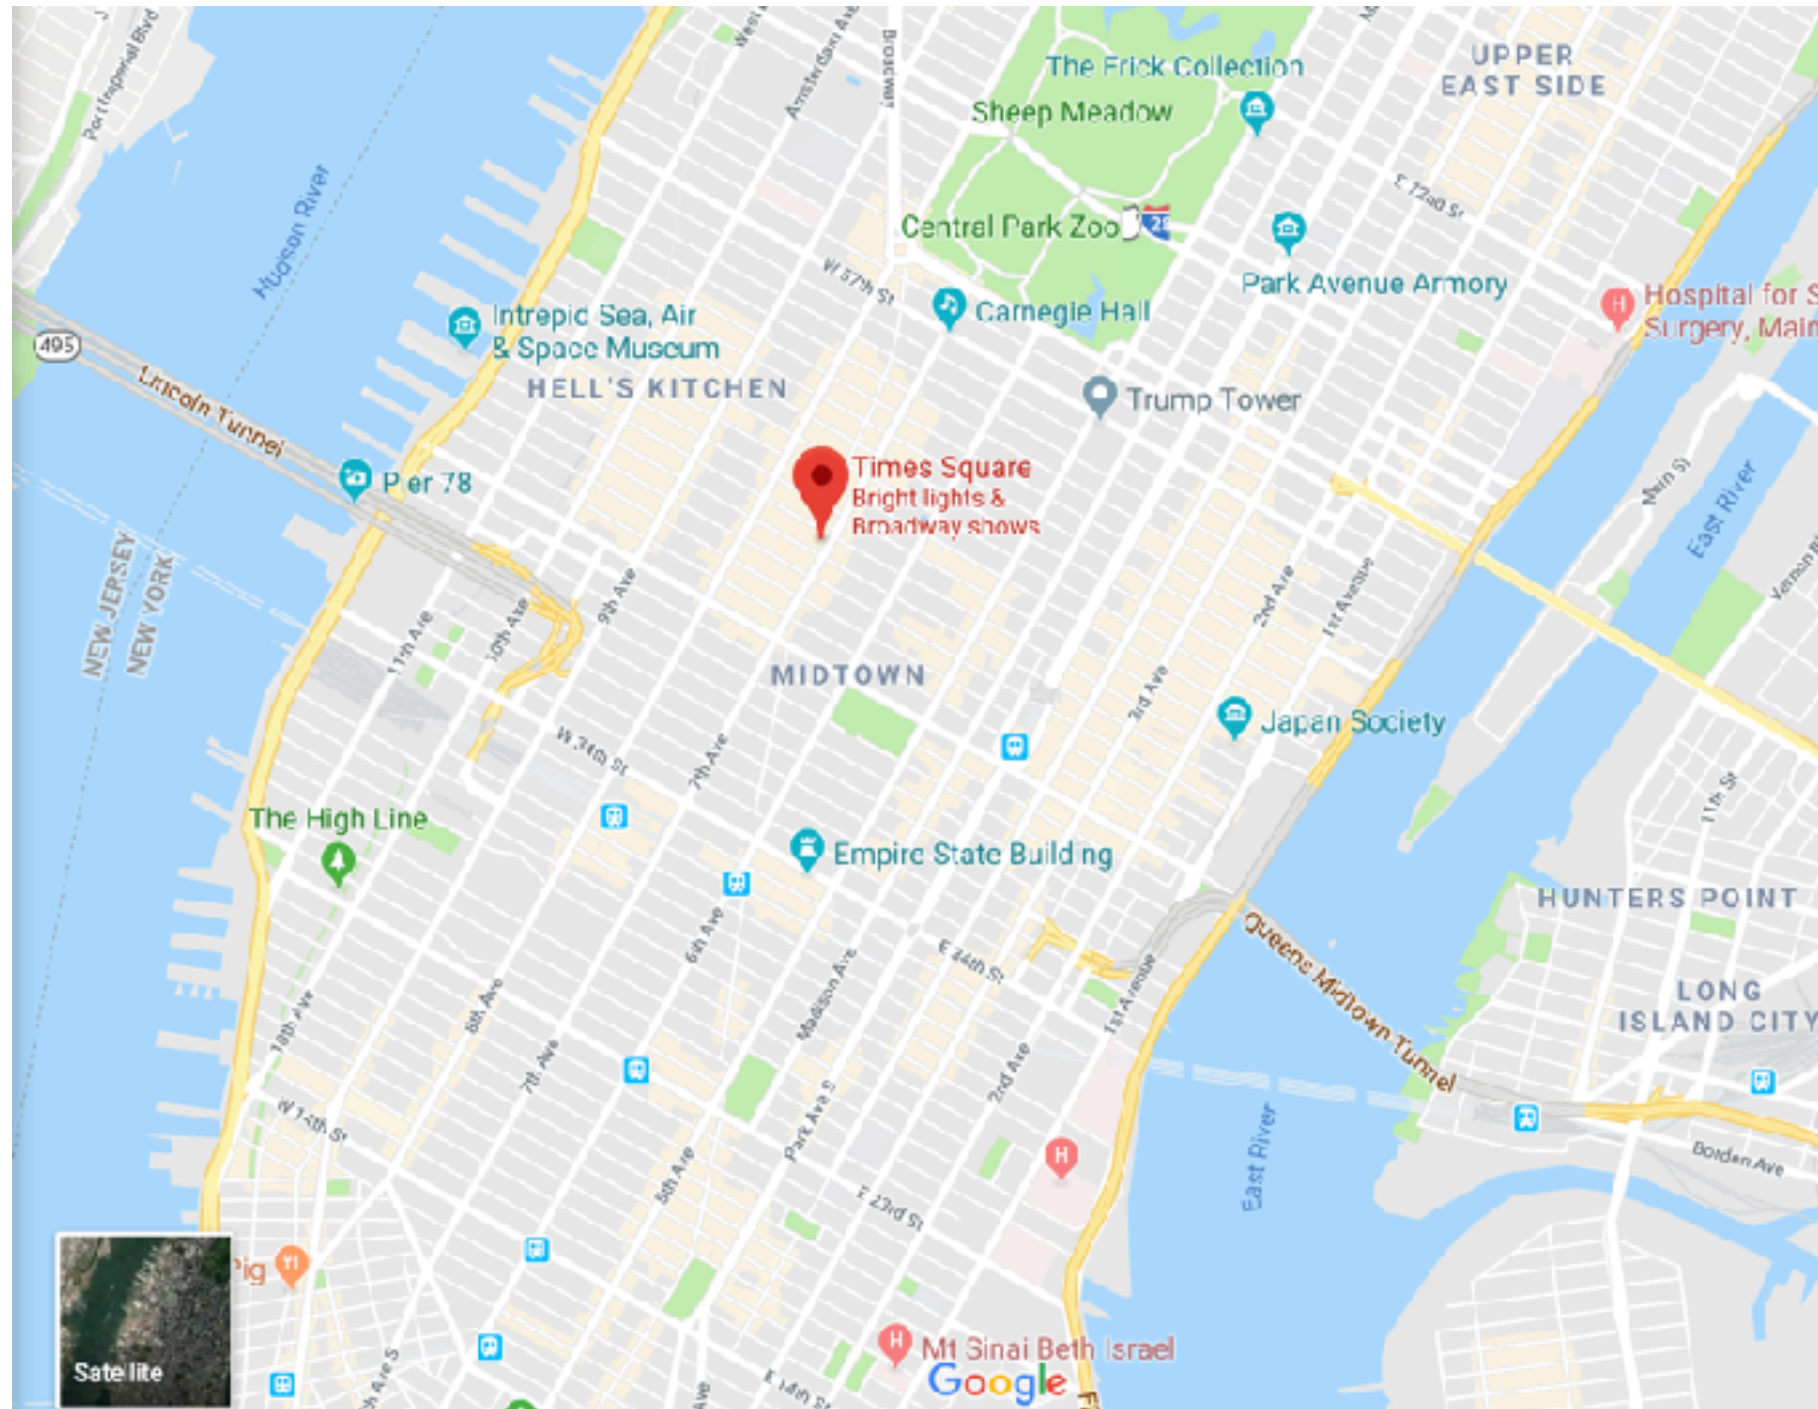


DISTANCE FUNCTION

- Generally, one popular heuristic function is a distance function.
- A distance function is a function to calculate a distance between two points in n dimensional space.
- Two most popular distance functions are
 - Manhattan distance
 - Euclidean distance

MANHATTAN DISTANCE

- The name is from a part of New York City: Manhattan.



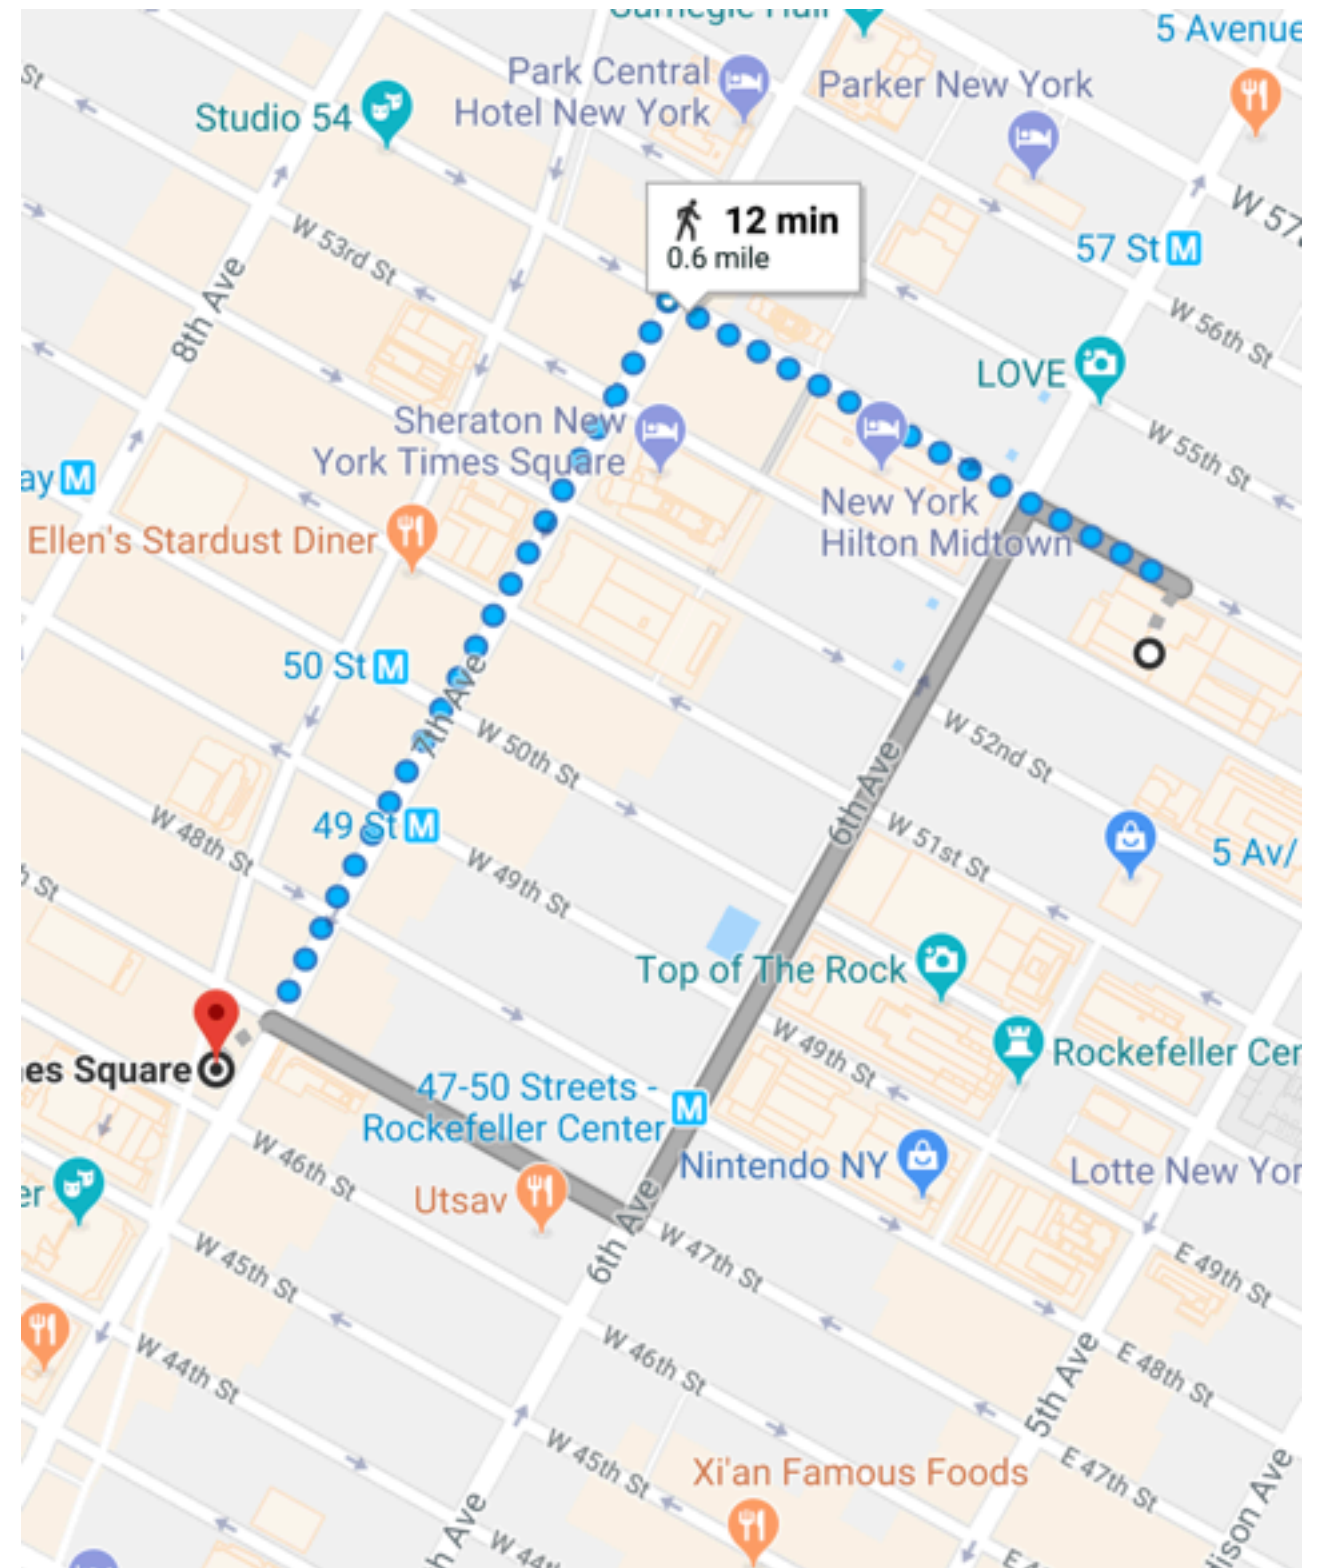
MANHATTAN DISTANCE

- The name is from a part of New York City: Manhattan.



MANHATTAN DISTANCE

- If we want to walk from one place to another, we need to walk along the street.
- And people count it as a “block”.
- Like, go north 5 for blocks, then east 3 for blocks.



MANHATTAN DISTANCE

➤ In 2-d:

$$d = \Delta x_1 + \Delta x_2$$

➤ In n-d:

$$d = \Delta x_1 + \Delta x_2 + \dots + \Delta x_n$$

EUCLIDEAN DISTANCE

- A fancy name of what we call “straight line distance”
- Can you guess what a straight line distance is?

EUCLIDEAN DISTANCE

► In 2-d:


$$d = \sqrt{\Delta x_1^2 + \Delta x_2^2}$$

► In n-d:

$$d = \sqrt{\Delta x_1^2 + \Delta x_2^2 + \dots + \Delta x_n^2}$$

MANHATTAN VS EUCLIDEAN

.....



Retrieved from : https://www.researchgate.net/profile/Katrina_Bolochko/publication/273487115/figure/fig3/AS:284032651808772@1444730069527/Fig-3-Manhattan-and-Euclidean-distances-between-two-points.png

GREEDY BEST-FIRST SEARCH

Informed Search

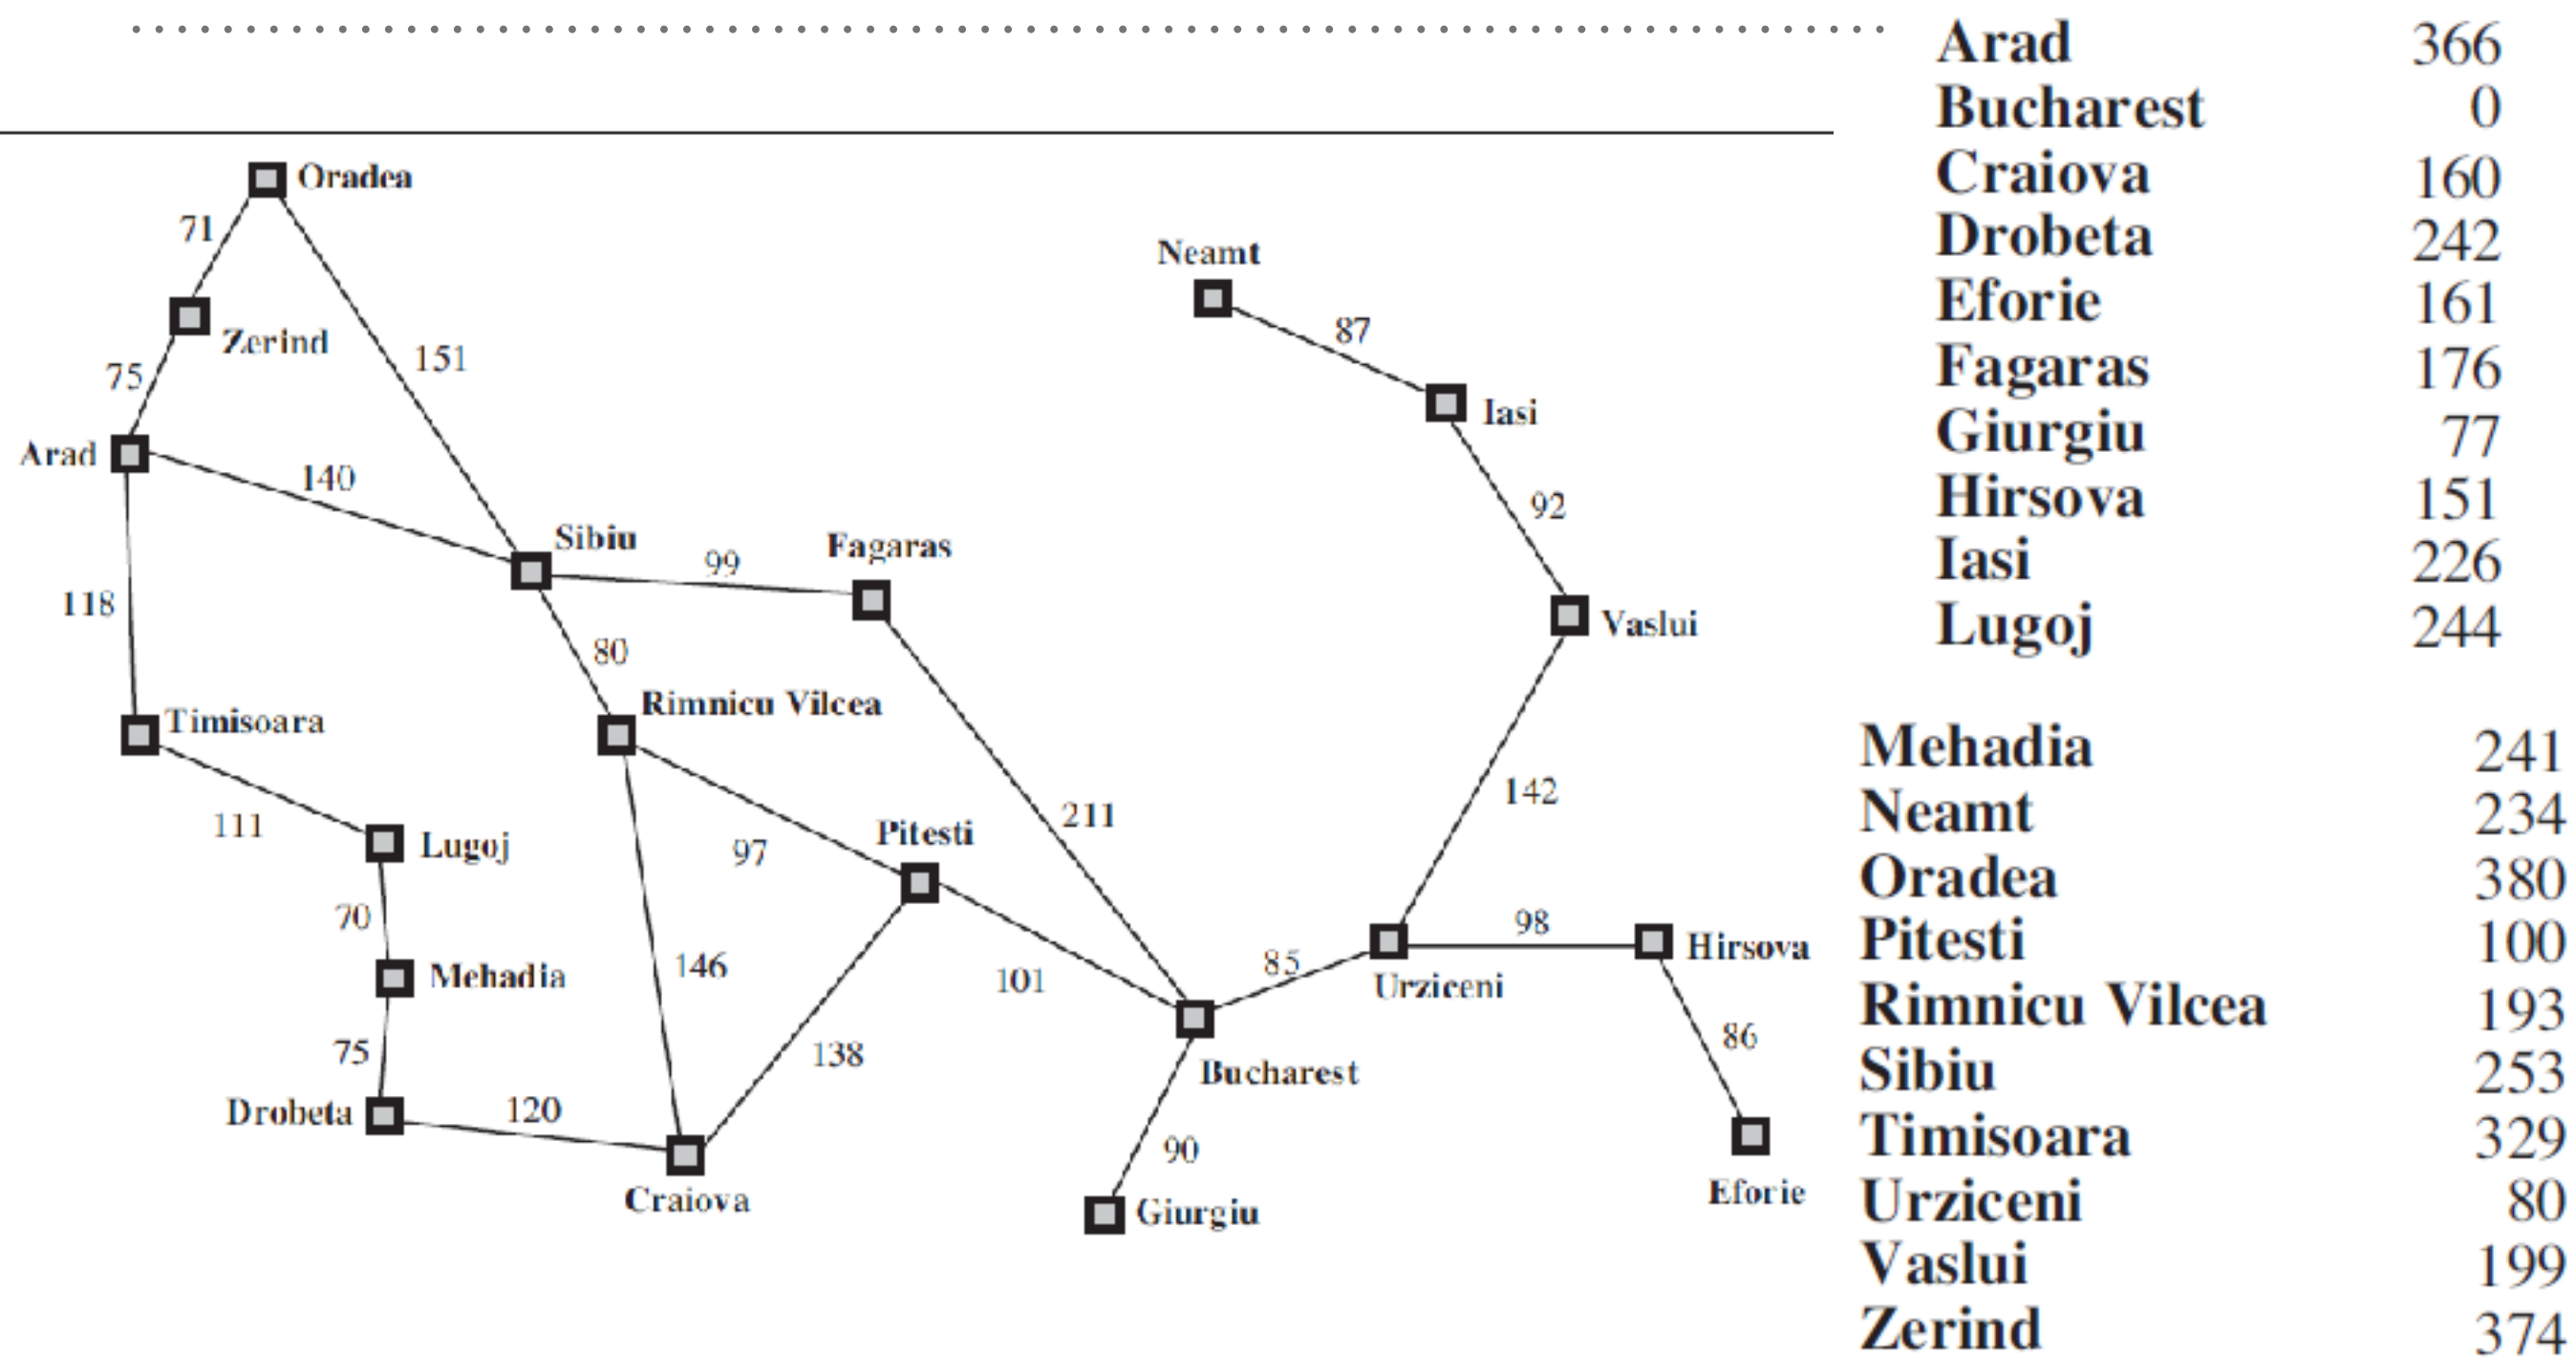
GREEDY BEST-FIRST SEARCH

Key idea:

- Expand the node that is the nearest to the goal first

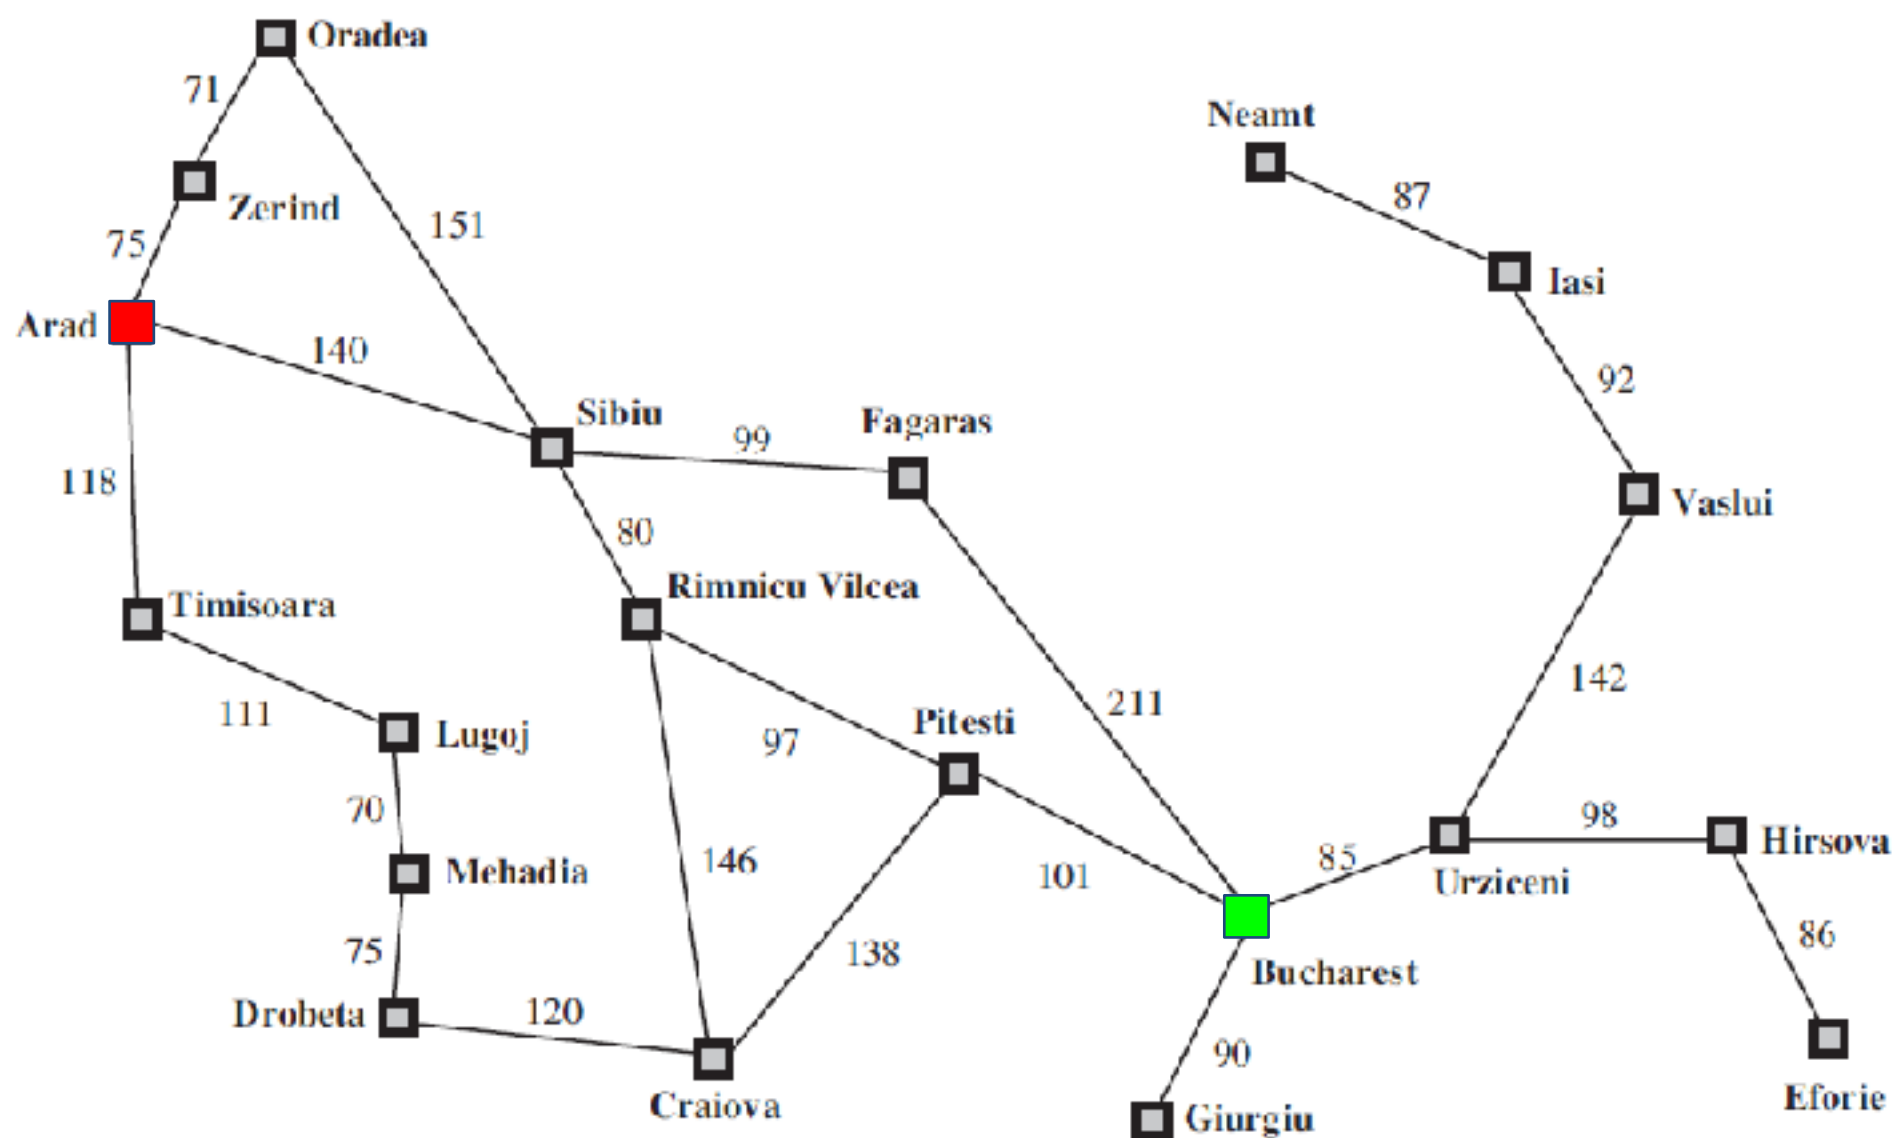
Definition:

- Greedy Best-first Search tries to expand the node that is closest to the goal, on the grounds that this is likely to lead to a solution quickly.

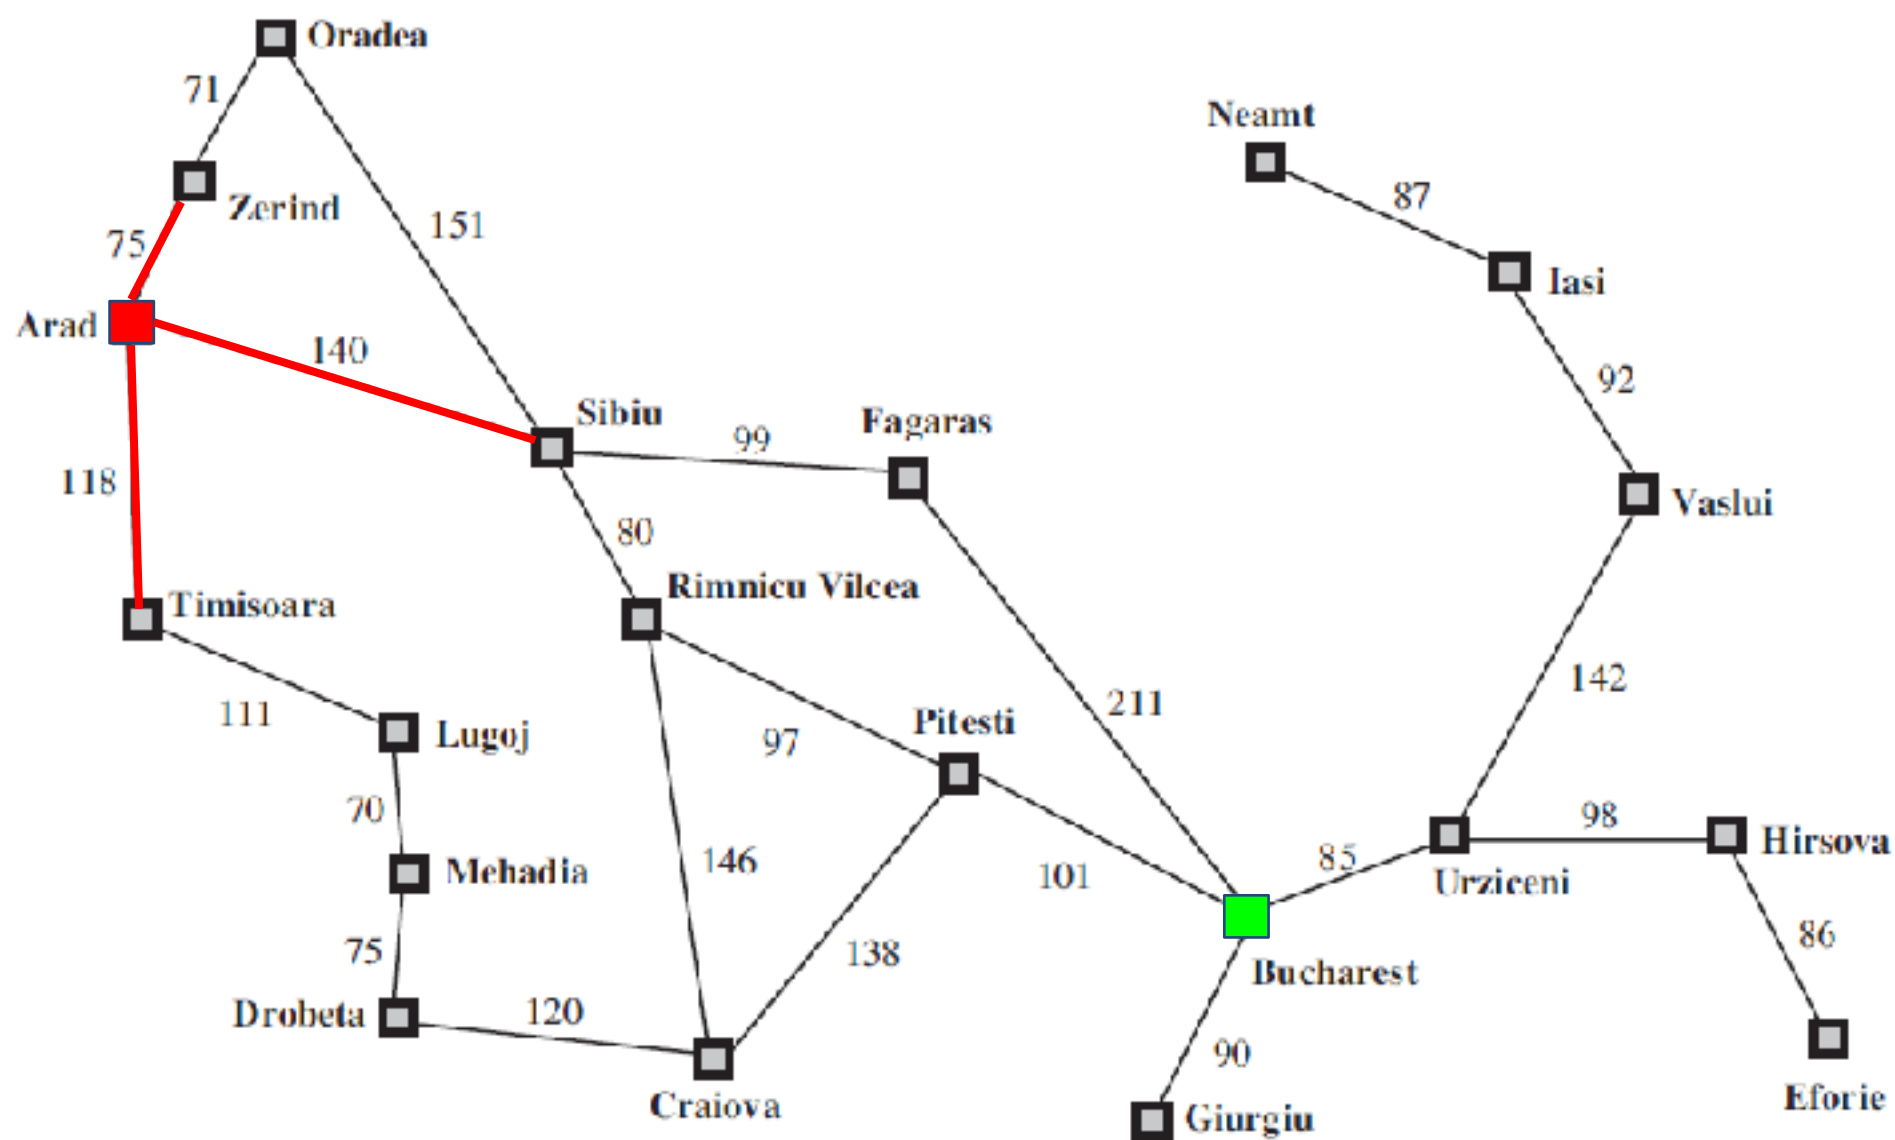


HEURISTIC FUNCTION


$h(n)$ = straight line distance from node n to the goal state.



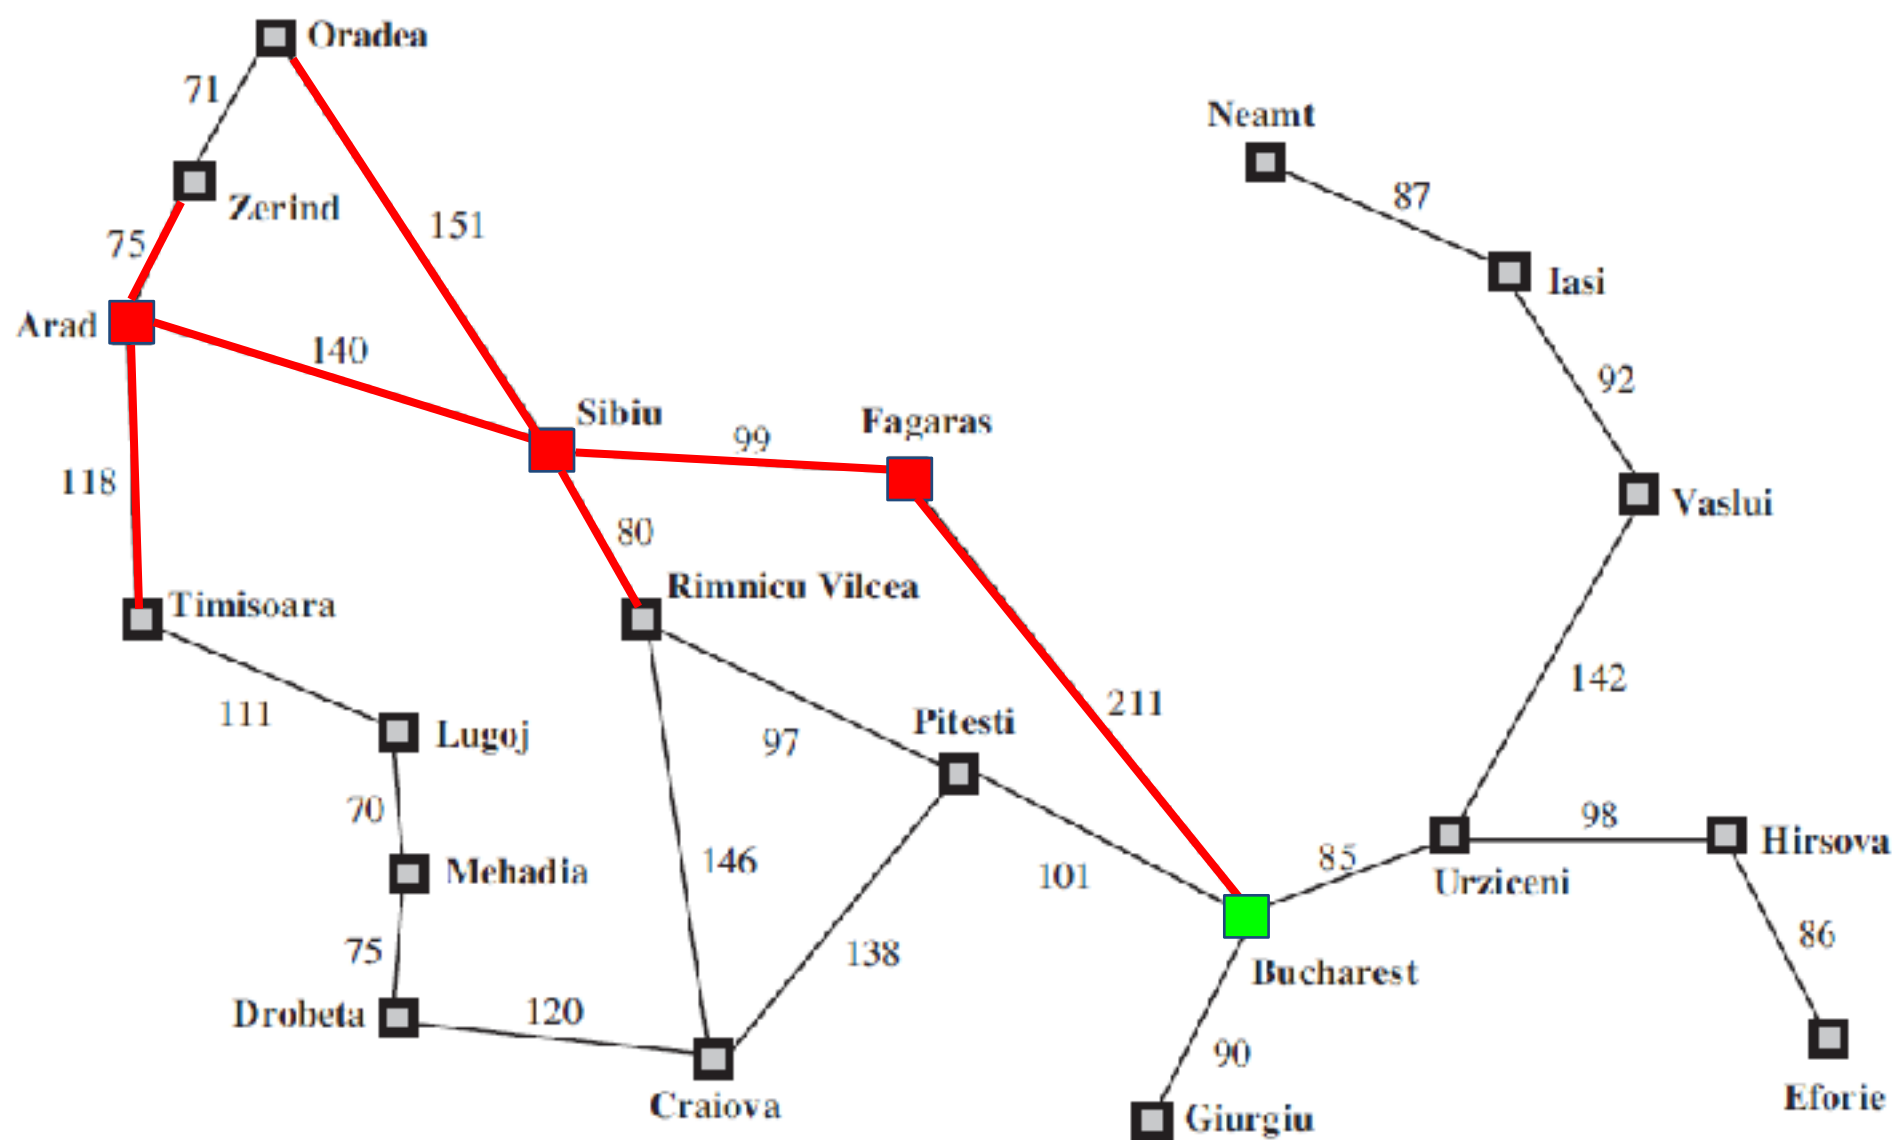
...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	Giurgiu	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	Oradea	380
	Pitesti	100
	Rimnicu Vilcea	193
	Sibiu	253
	Timisoara	329
	Urziceni	80
	Vaslui	199
	Zerind	374



...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	Giurgiu	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	Oradea	380
	Pitesti	100
	Rimnicu Vilcea	193
	<u>Sibiu</u>	<u>253</u>
	<u>Timisoara</u>	<u>329</u>
	Urziceni	80
	Vaslui	199
	<u>Zerind</u>	<u>374</u>



...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	<u>Giurgiu</u>	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	<u>Oradea</u>	380
	Pitesti	100
	<u>Rimnicu Vilcea</u>	193
	Sibiu	253
	Timisoara	329
	Urziceni	80
	Vaslui	199
	Zerind	374



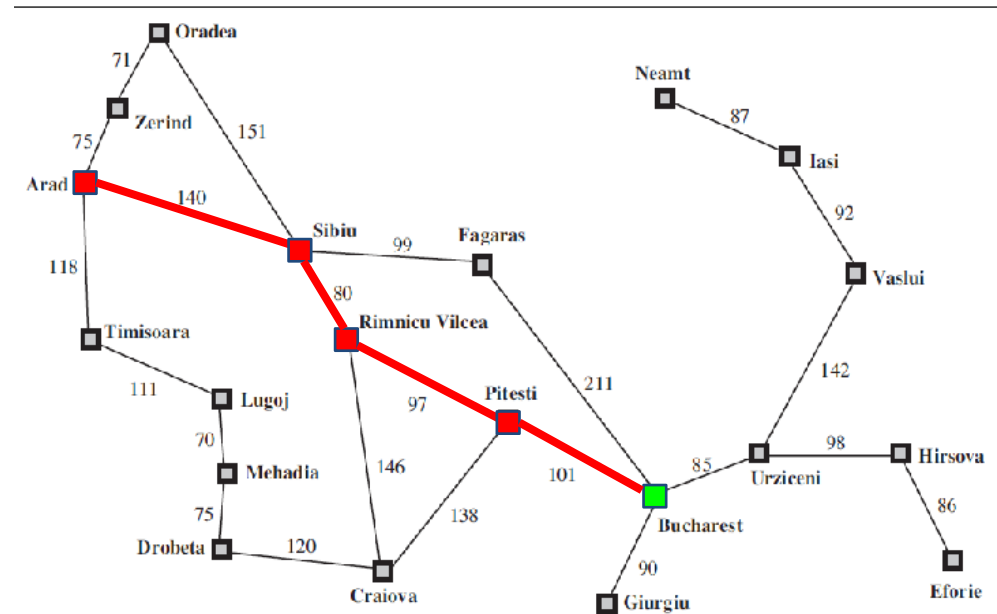

...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	<u>Giurgiu</u>	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	<u>Oradea</u>	380
	Pitesti	100
	<u>Rimnicu Vilcea</u>	193
	Sibiu	253
	Timisoara	329
	Urziceni	80
	Vaslui	199
	Zerind	374

GREEDY BEST-FIRST SEARCH

.....

It looks good, but it is not optimal.

- Path on the left is 32 kilometers longer.



A* SEARCH

Informed Search


A* SEARCH (A-STAR SEARCH)

- Key Idea:
 - Expands the cheapest node
 - Cheapest node = minimum cost of (initial state to node n + node n to goal state)
- Definition:
 - A* Search is a form of best-first search. It evaluates nodes by combining $g(n)$, the cost to reach the node, and $h(n)$ the cost to get from the node to the goal:
 - $f(n) = g(n) + h(n)$


EVALUATION FUNCTION

- A* Search is a form of best-first search. It evaluates nodes by combining $g(n)$, the cost to reach the node, and $h(n)$ the cost to get from the node to the goal:
- $f(n) = g(n) + h(n)$
 - $f(n)$ = Evaluation function
 - $g(n)$ = Cost function from initial state to current state
 - $h(n)$ = Heuristic function from current state to goal state


(a) The initial state

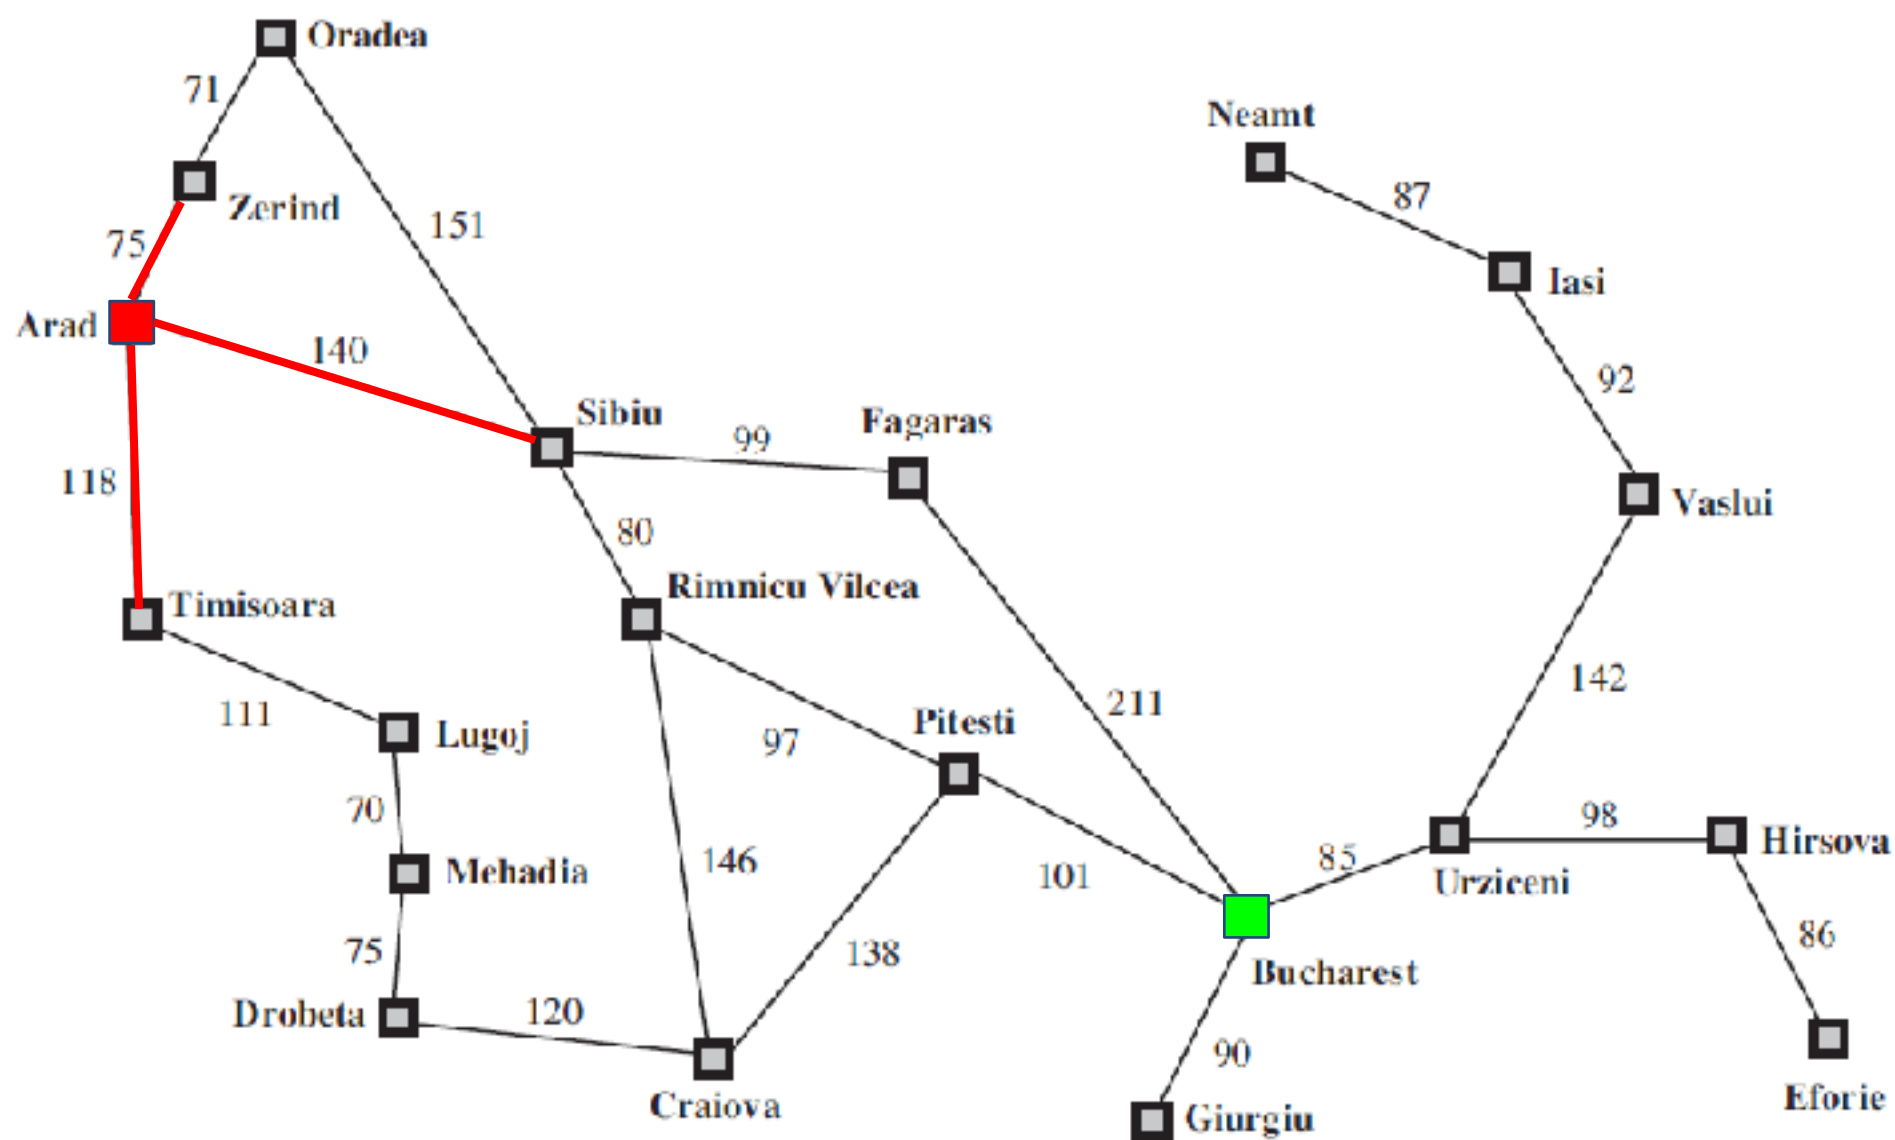


(b) After expanding Arad




(c) After expanding Sibiu






...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	Giurgiu	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	Oradea	380
	Pitesti	100
	Rimnicu Vilcea	193
	Sibiu	253
	Timisoara	329
	Urziceni	80
	Vaslui	199
	Zerind	374


(a) The initial state

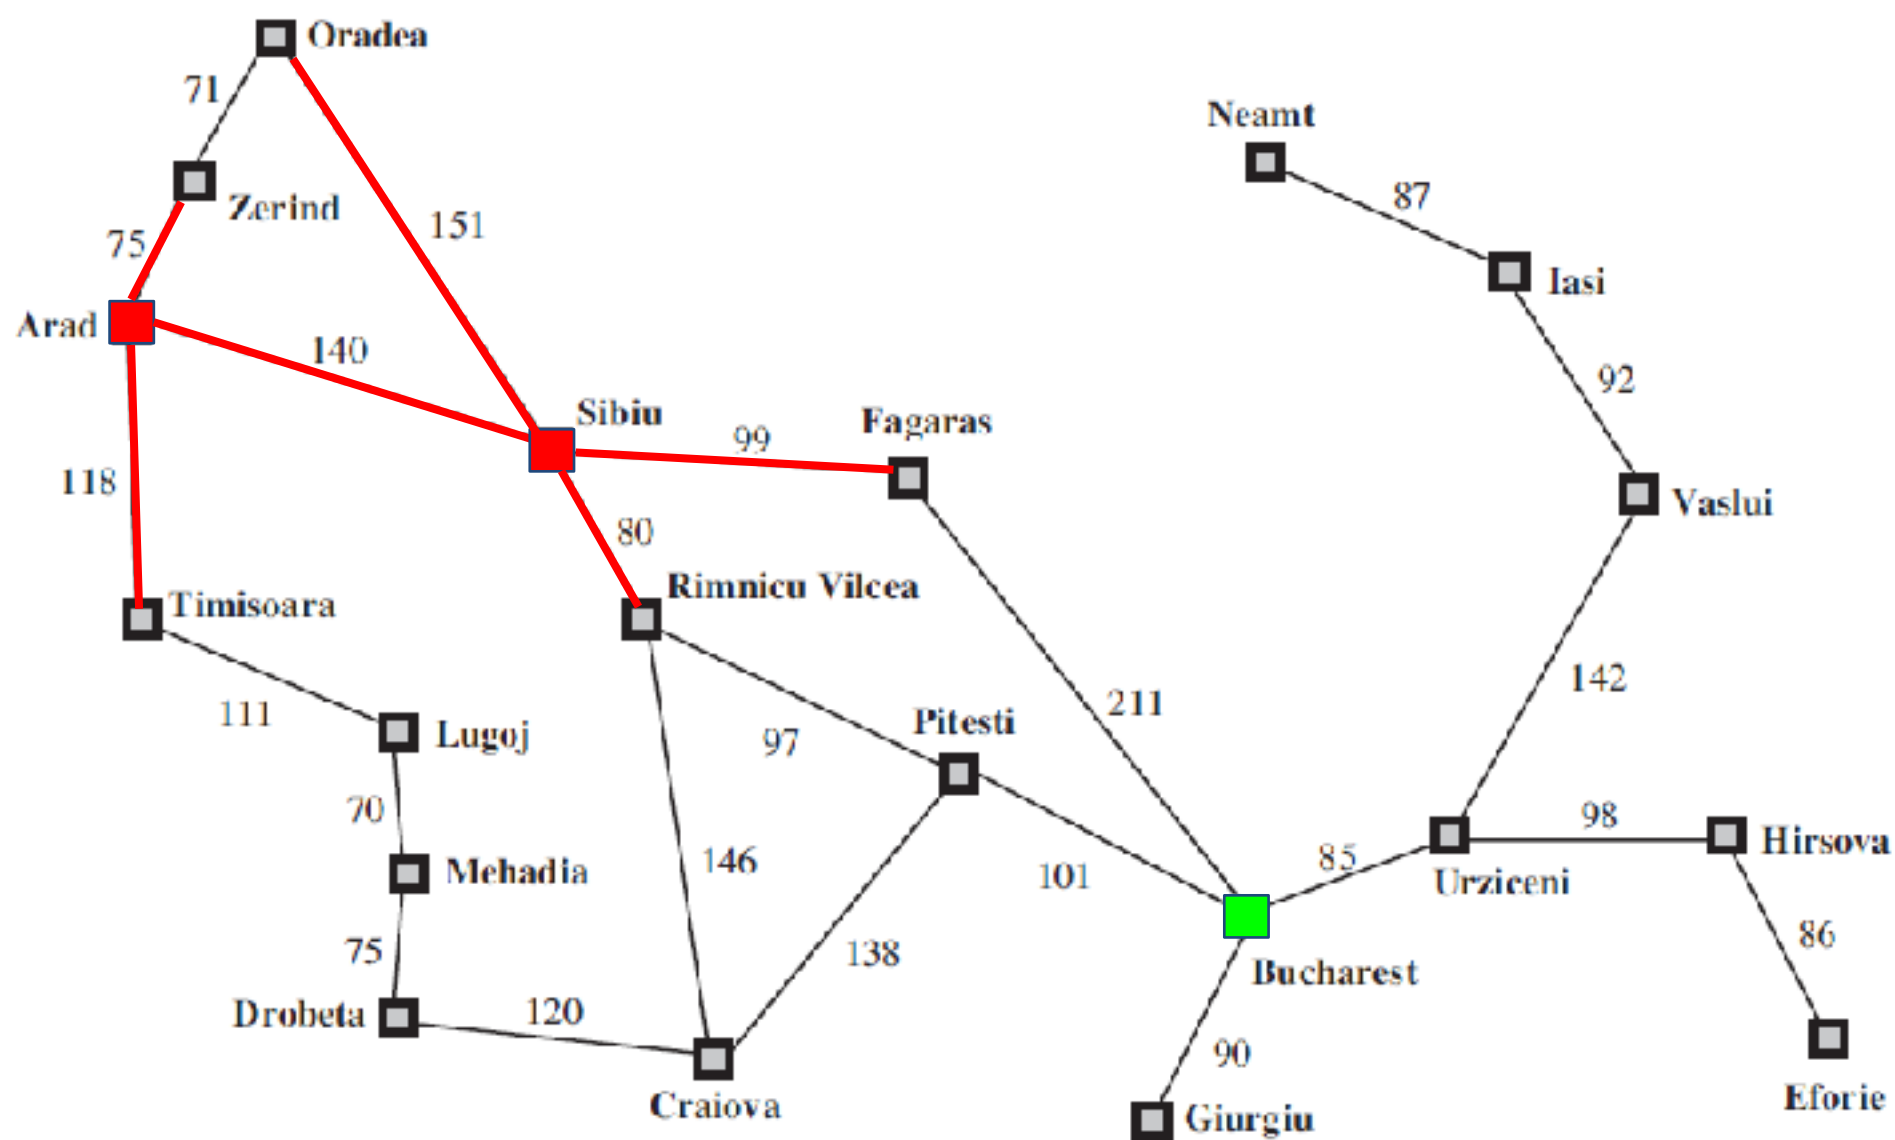


(b) After expanding Arad




(c) After expanding Sibiu






...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	Giurgiu	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	Oradea	380
	Pitesti	100
	Rimnicu Vilcea	193
	Sibiu	253
	Timisoara	329
	Urziceni	80
	Vaslui	199
	Zerind	374


(a) The initial state




(b) After expanding Arad



(c) After expanding Sibiu

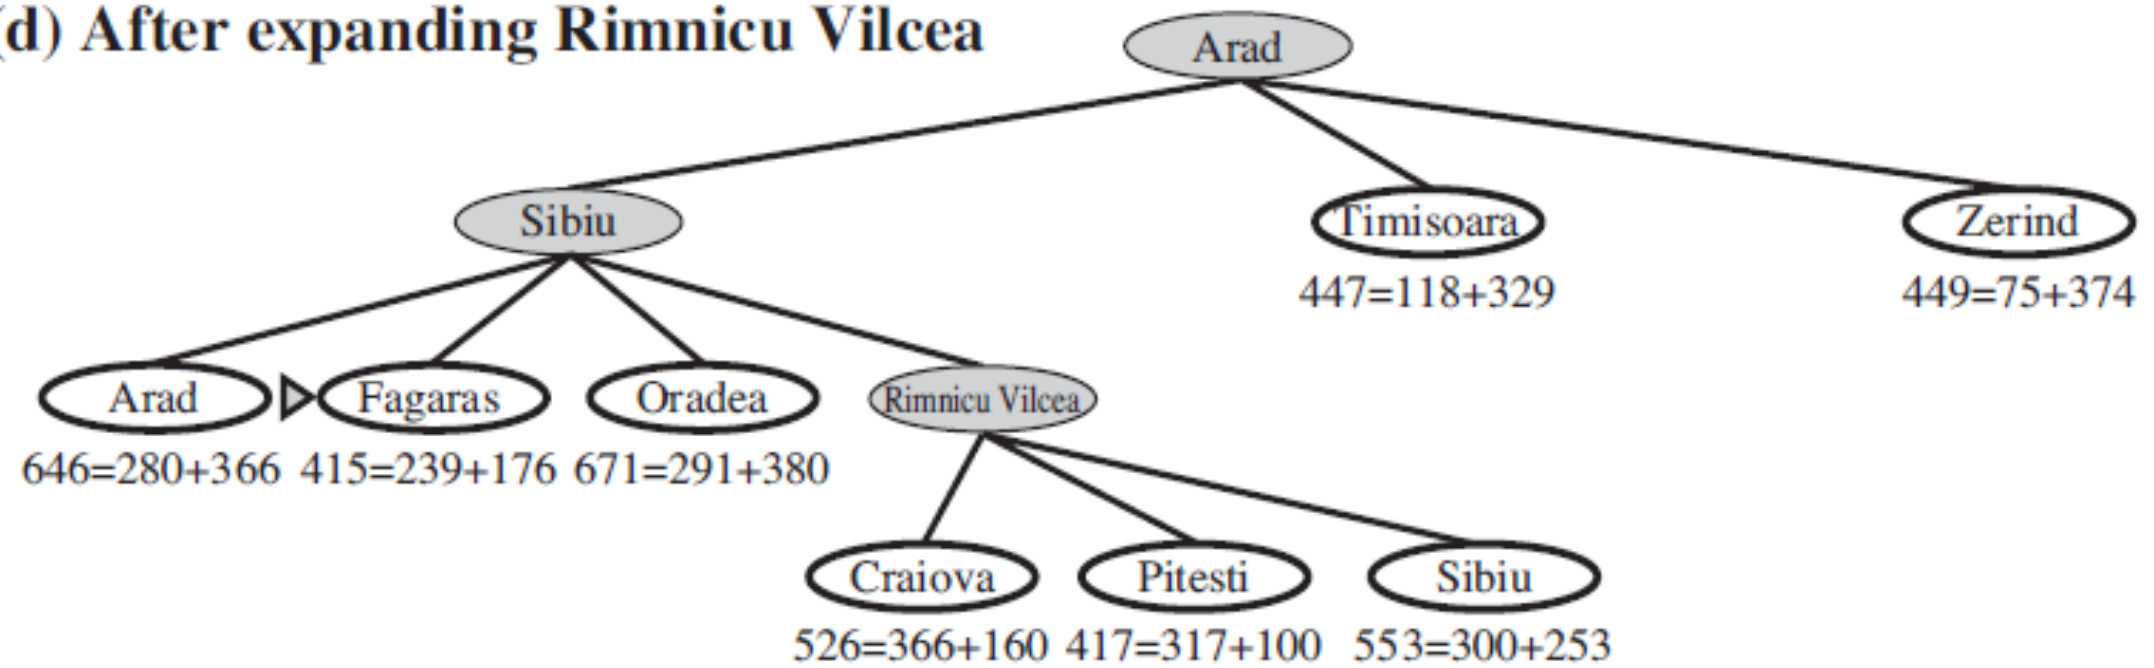


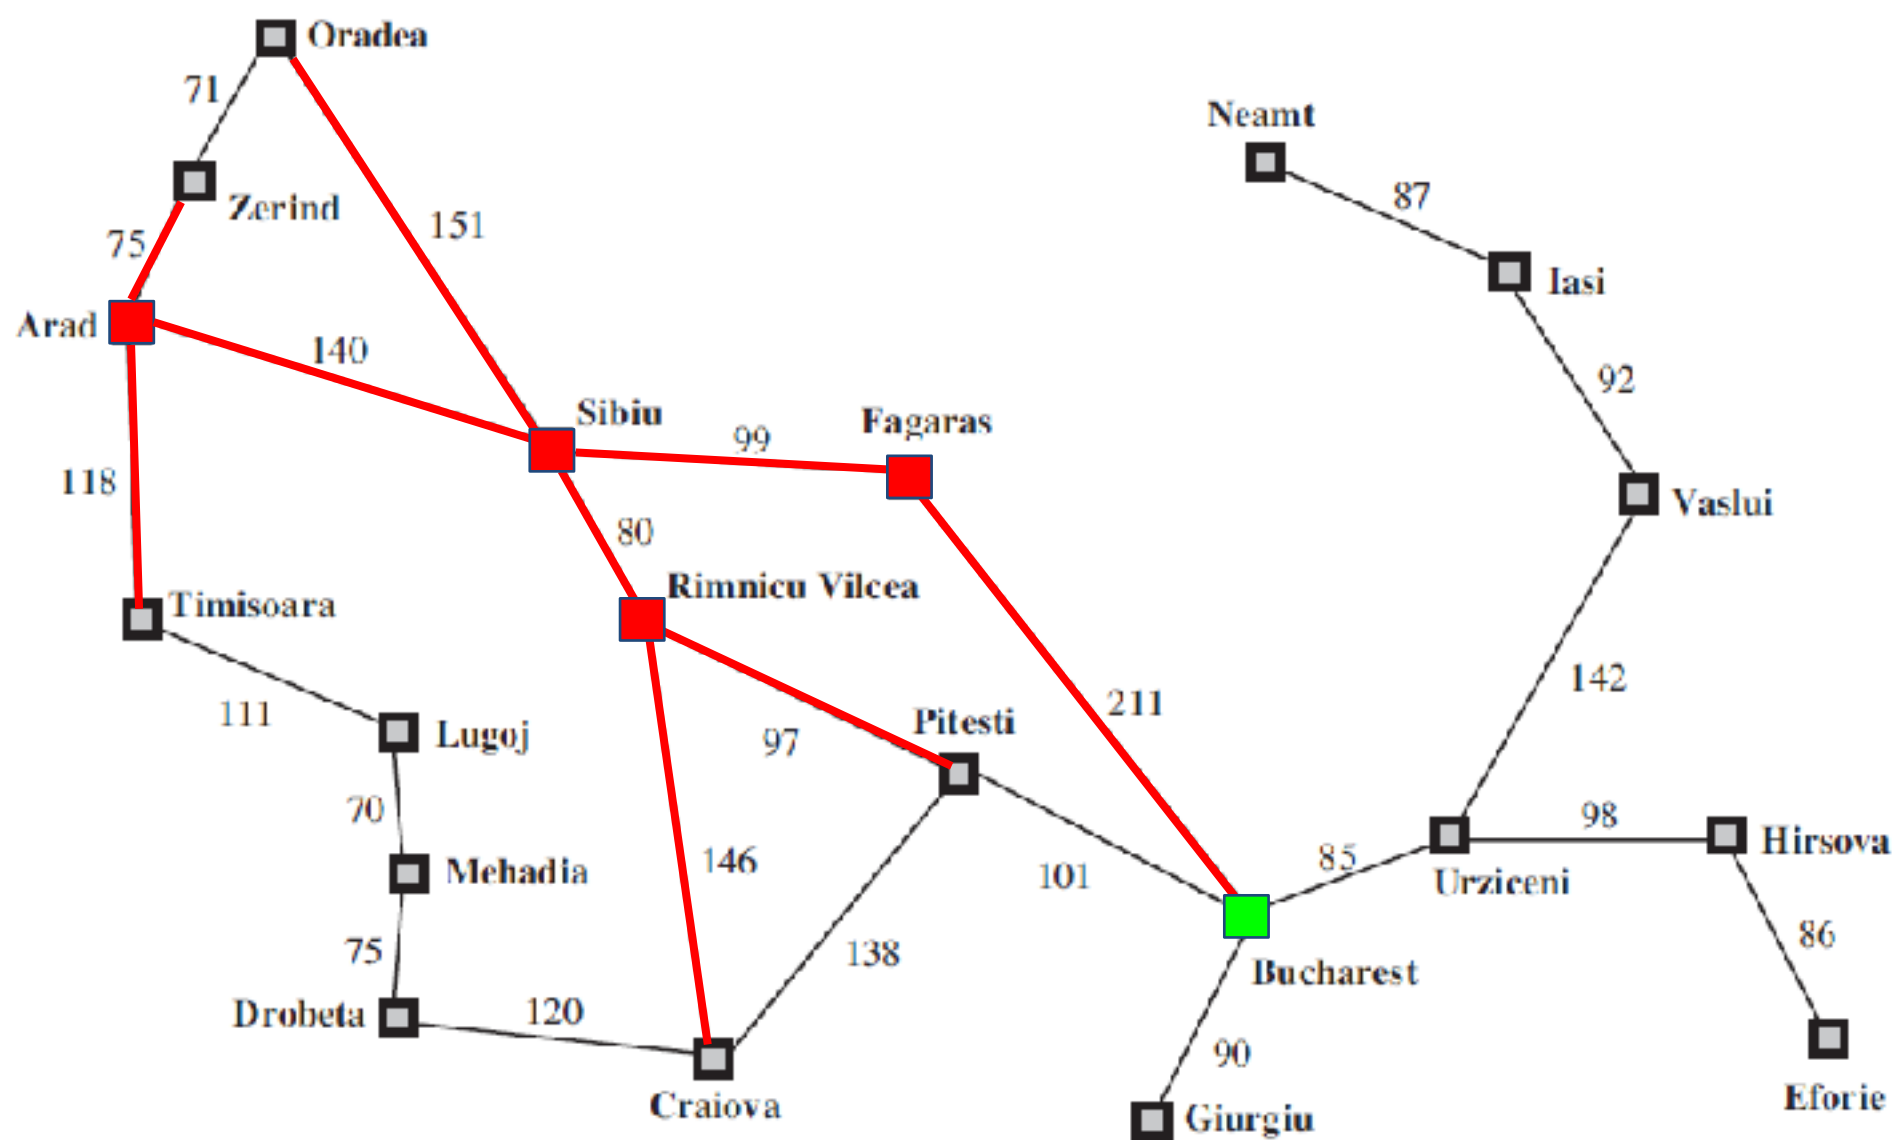


...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	Giurgiu	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	Oradea	380
	Pitesti	100
	Rimnicu Vilcea	193
	Sibiu	253
	Timisoara	329
	Urziceni	80
	Vaslui	199
	Zerind	374

.....

(d) After expanding Rimnicu Vilcea

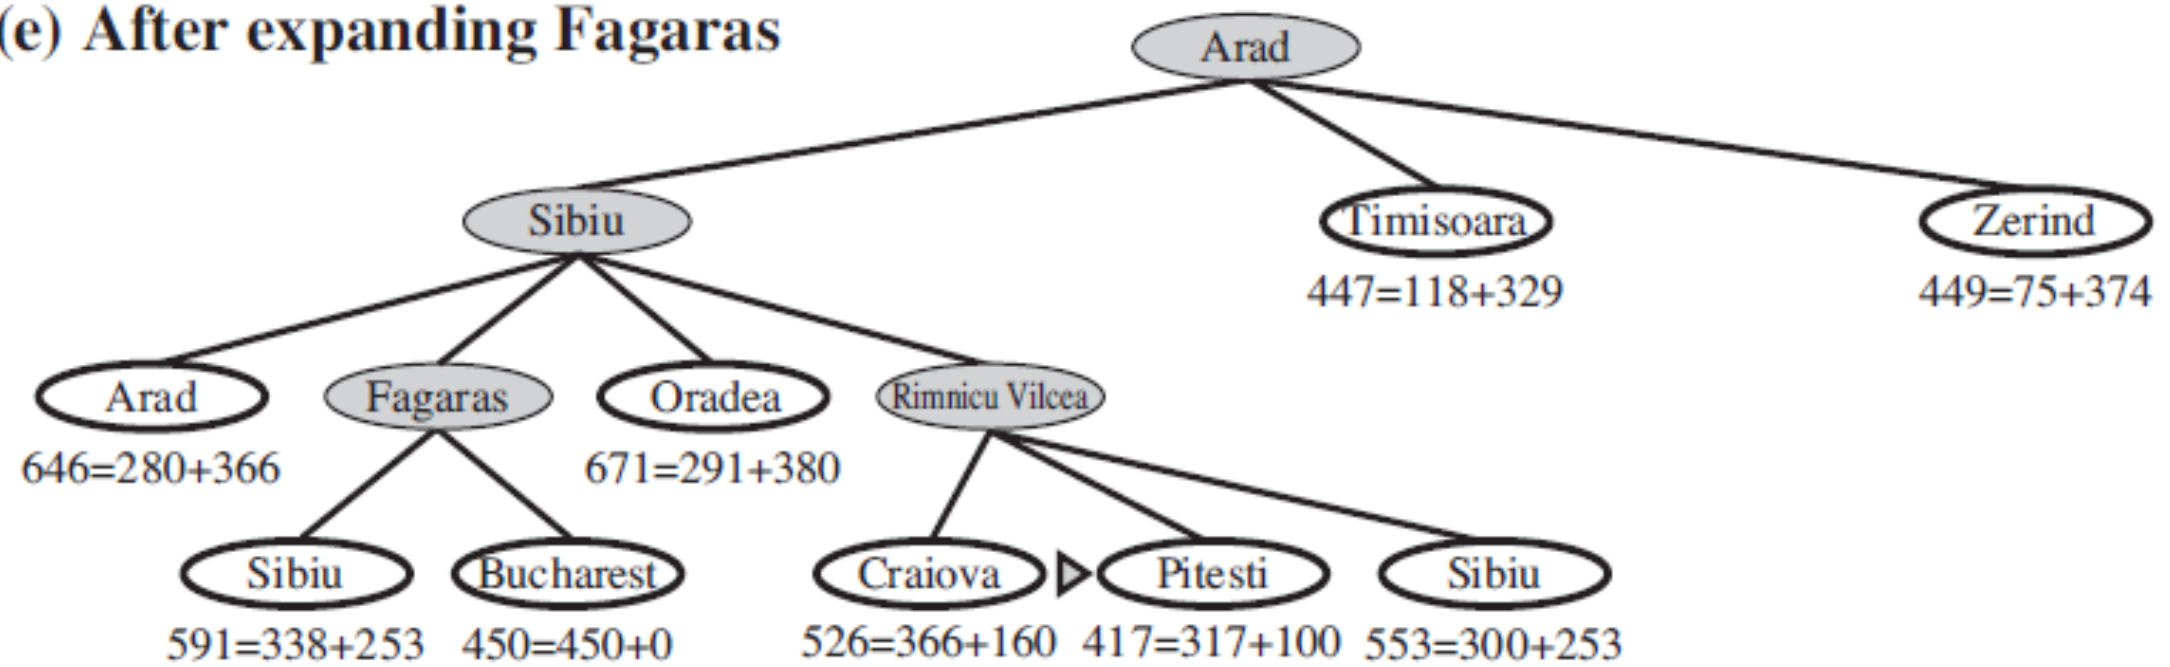


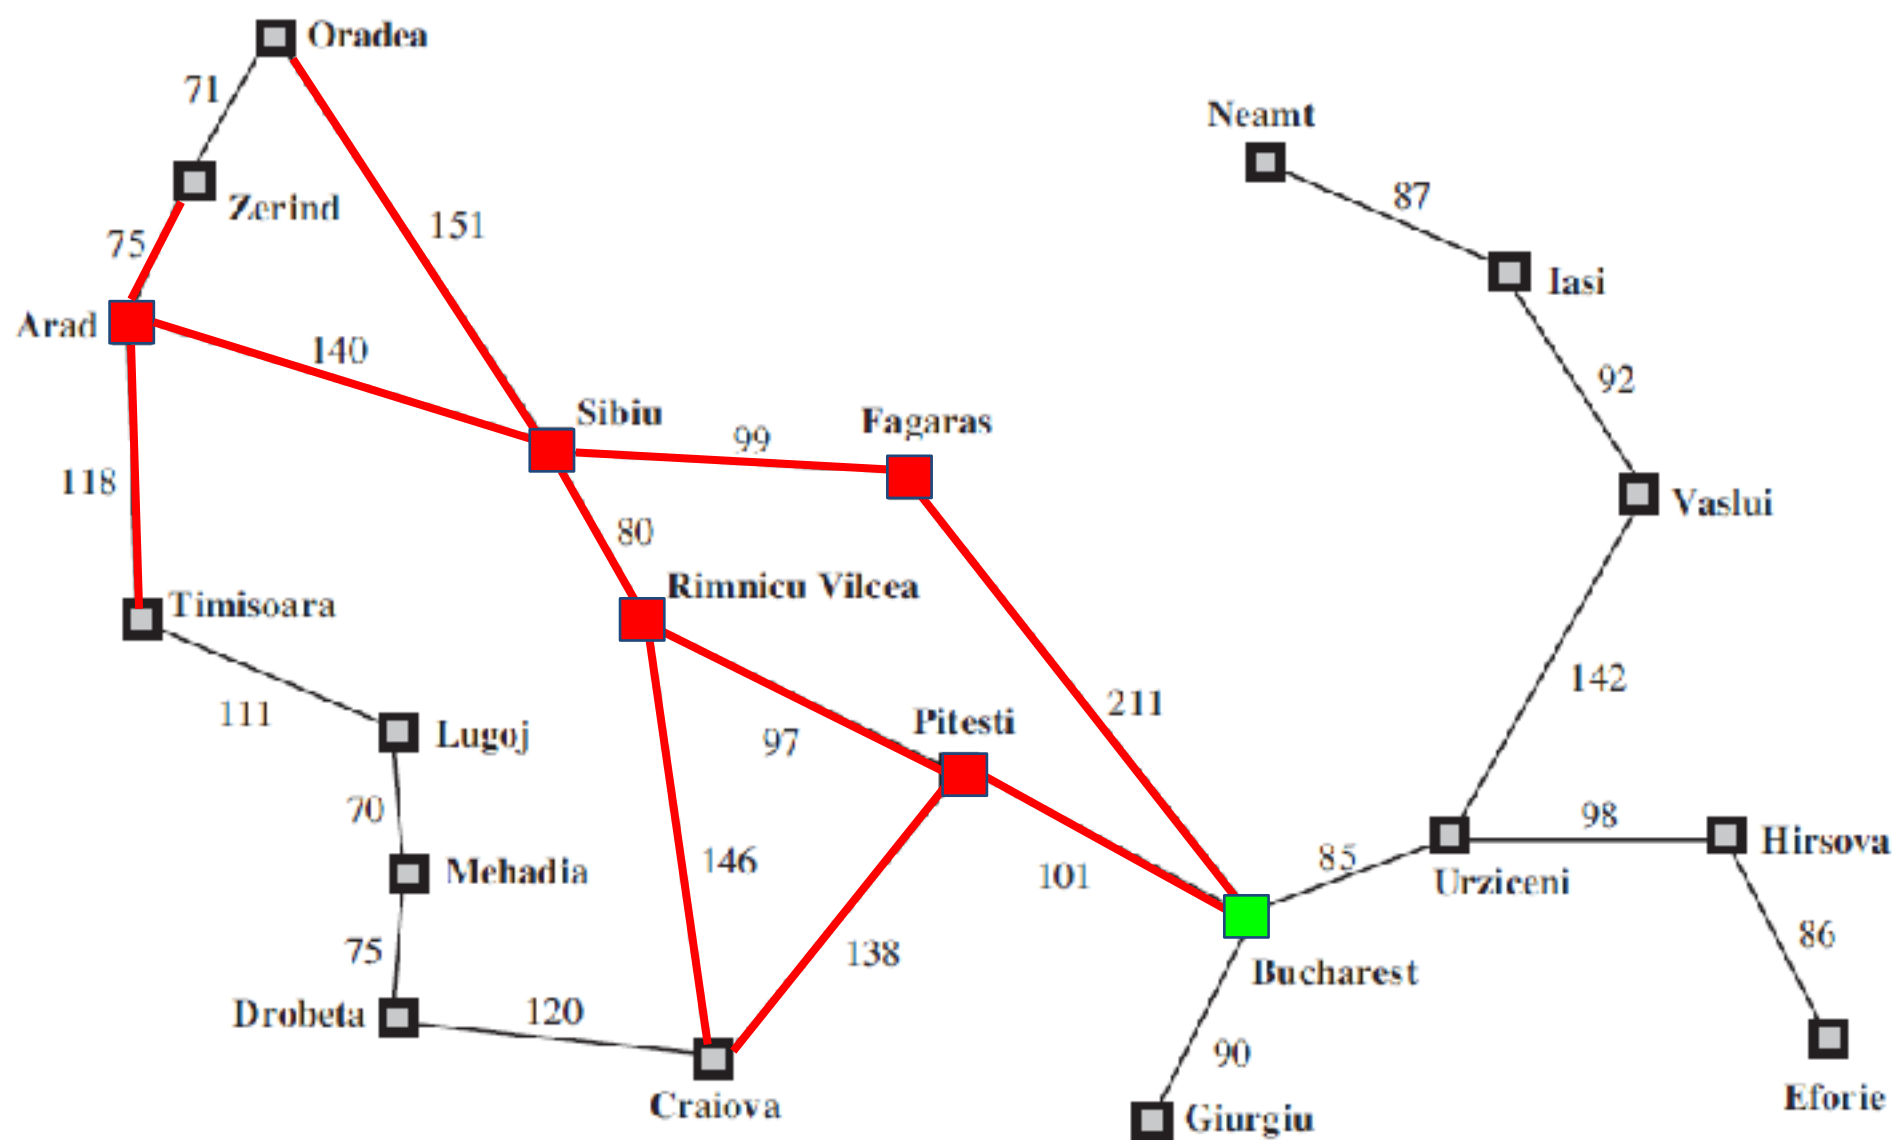


...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	Giurgiu	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	Oradea	380
	Pitesti	100
	Rimnicu Vilcea	193
	Sibiu	253
	Timisoara	329
	Urziceni	80
	Vaslui	199
	Zerind	374

.....

(e) After expanding Fagaras

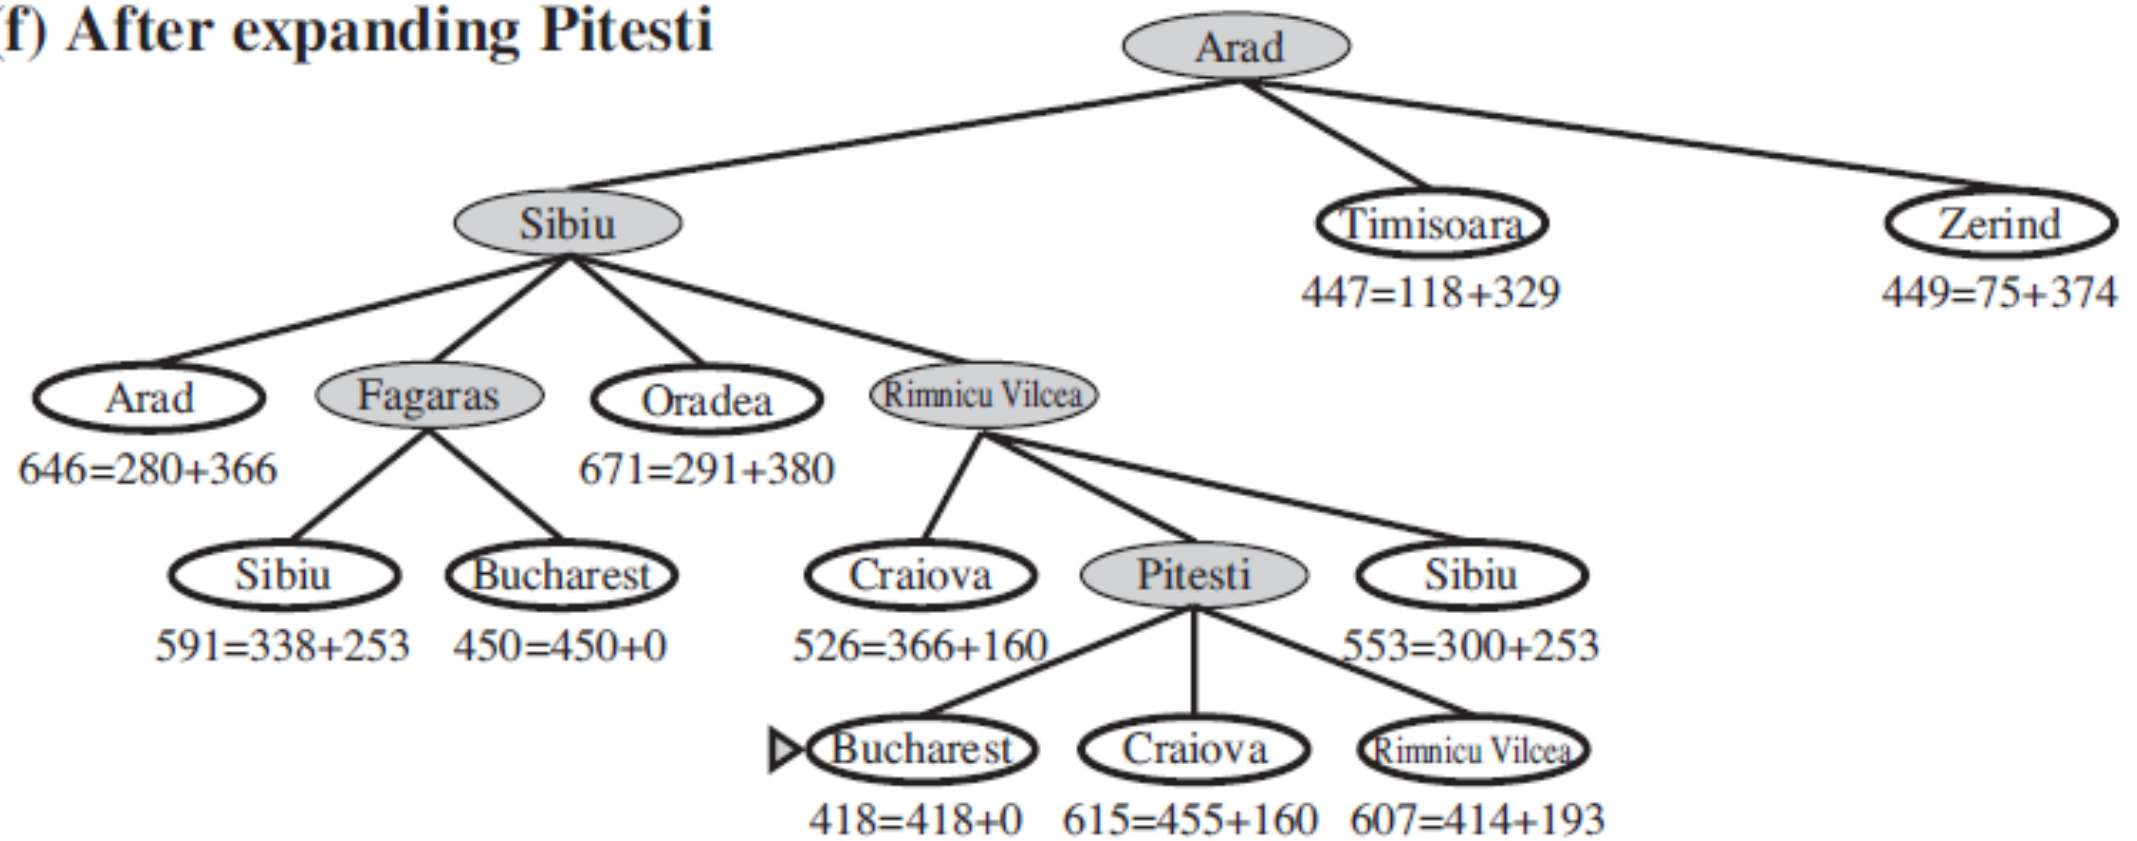




...	Arad	366
	Bucharest	0
	Craiova	160
	Drobeta	242
	Eforie	161
	Fagaras	176
	Giurgiu	77
	Hirsova	151
	Iasi	226
	Lugoj	244
	Mehadia	241
	Neamt	234
	Oradea	380
	Pitesti	100
	Rimnicu Vilcea	193
	Sibiu	253
	Timisoara	329
	Urziceni	80
	Vaslui	199
	Zerind	374

.....

(f) After expanding Pitesti



A* QUIZ

.....

	1	2	3	4	5
1	s				
2					
3					
4					
5		g			

Starting state = (1,1)

Goal state = (2,5)

g. Heuristic function ที่จะเลือกใช้แก้ปัญหาแผนที่นี้ จะเลือกใช้อะไร และสิ่งที่เลือกใช้ คืออะไร

h. Evaluation function สำหรับการหา A* Search บนแผนที่นี้ คืออะไร แต่ละองค์ประกอบของ evaluation function หมายความว่าอะไรบ้าง

i. จงวาด tree ที่เกิดจากการหา path planning โดยใช้ A* Algorithm ให้ใส่ตัวเลขที่แต่ละกิ่ง(branch) เพื่อแสดงลำดับการแตกกิ่ง

THE ALGORITHM

.....

function RECURSIVE-BEST-FIRST-SEARCH(*problem*) **returns** a solution, or failure
 return RBFS(*problem*, MAKE-NODE(*problem*.INITIAL-STATE), ∞)

function RBFS(*problem*, *node*, *f-limit*) **returns** a solution, or failure and a new *f*-cost limit
 if *problem*.GOAL-TEST(*node*.STATE) **then return** SOLUTION(*node*)
 successors $\leftarrow []$
 for each *action* **in** *problem*.ACTIONS(*node*.STATE) **do**
 add CHILD-NODE(*problem*, *node*, *action*) into *successors*
 if *successors* is empty **then return** failure, ∞
 for each *s* **in** *successors* **do** /* update *f* with value from previous search, if any */
 s.f $\leftarrow \max(s.g + s.h, \text{node.f})$
 loop do
 best \leftarrow the lowest *f*-value node in *successors*
 if *best.f* > *f-limit* **then return** failure, *best.f*
 alternative \leftarrow the second-lowest *f*-value among *successors*
 result, *best.f* \leftarrow RBFS(*problem*, *best*, min(*f-limit*, *alternative*))
 if *result* \neq failure **then return** *result*

THE ALGORITHM

- Another way to implement A^*

https://en.wikipedia.org/wiki/A*_search_algorithm

ADVANCED TOPICS

*Nice to know, but not a must-know
(We will hold this lecture later)*

LOCAL SEARCH & OPTIMIZATION PROBLEM


LOCAL SEARCH

- Key Idea:

- Have only one node/robot. Try to find the goal.

- Definition:

- The algorithms operate using a single current node and generally move only to neighbors of that node.



HILL CLIMBING SEARCH

Local Search and Optimization Problem

HILL CLIMBING SEARCH


- Key Idea:

- Just like climbing the hill, trying to reach the top
- Always go uphill (higher)

- Definition:

- The algorithm is simply a loop that continually moves in the direction of increasing value - that is, uphill. It terminates when it reaches a “peak” where no neighbor has a higher value.

EXAMPLE



HILL CLIMBING: THE ALGORITHM

function HILL-CLIMBING(*problem*) **returns** a state that is a local maximum

current \leftarrow MAKE-NODE(*problem*.INITIAL-STATE)

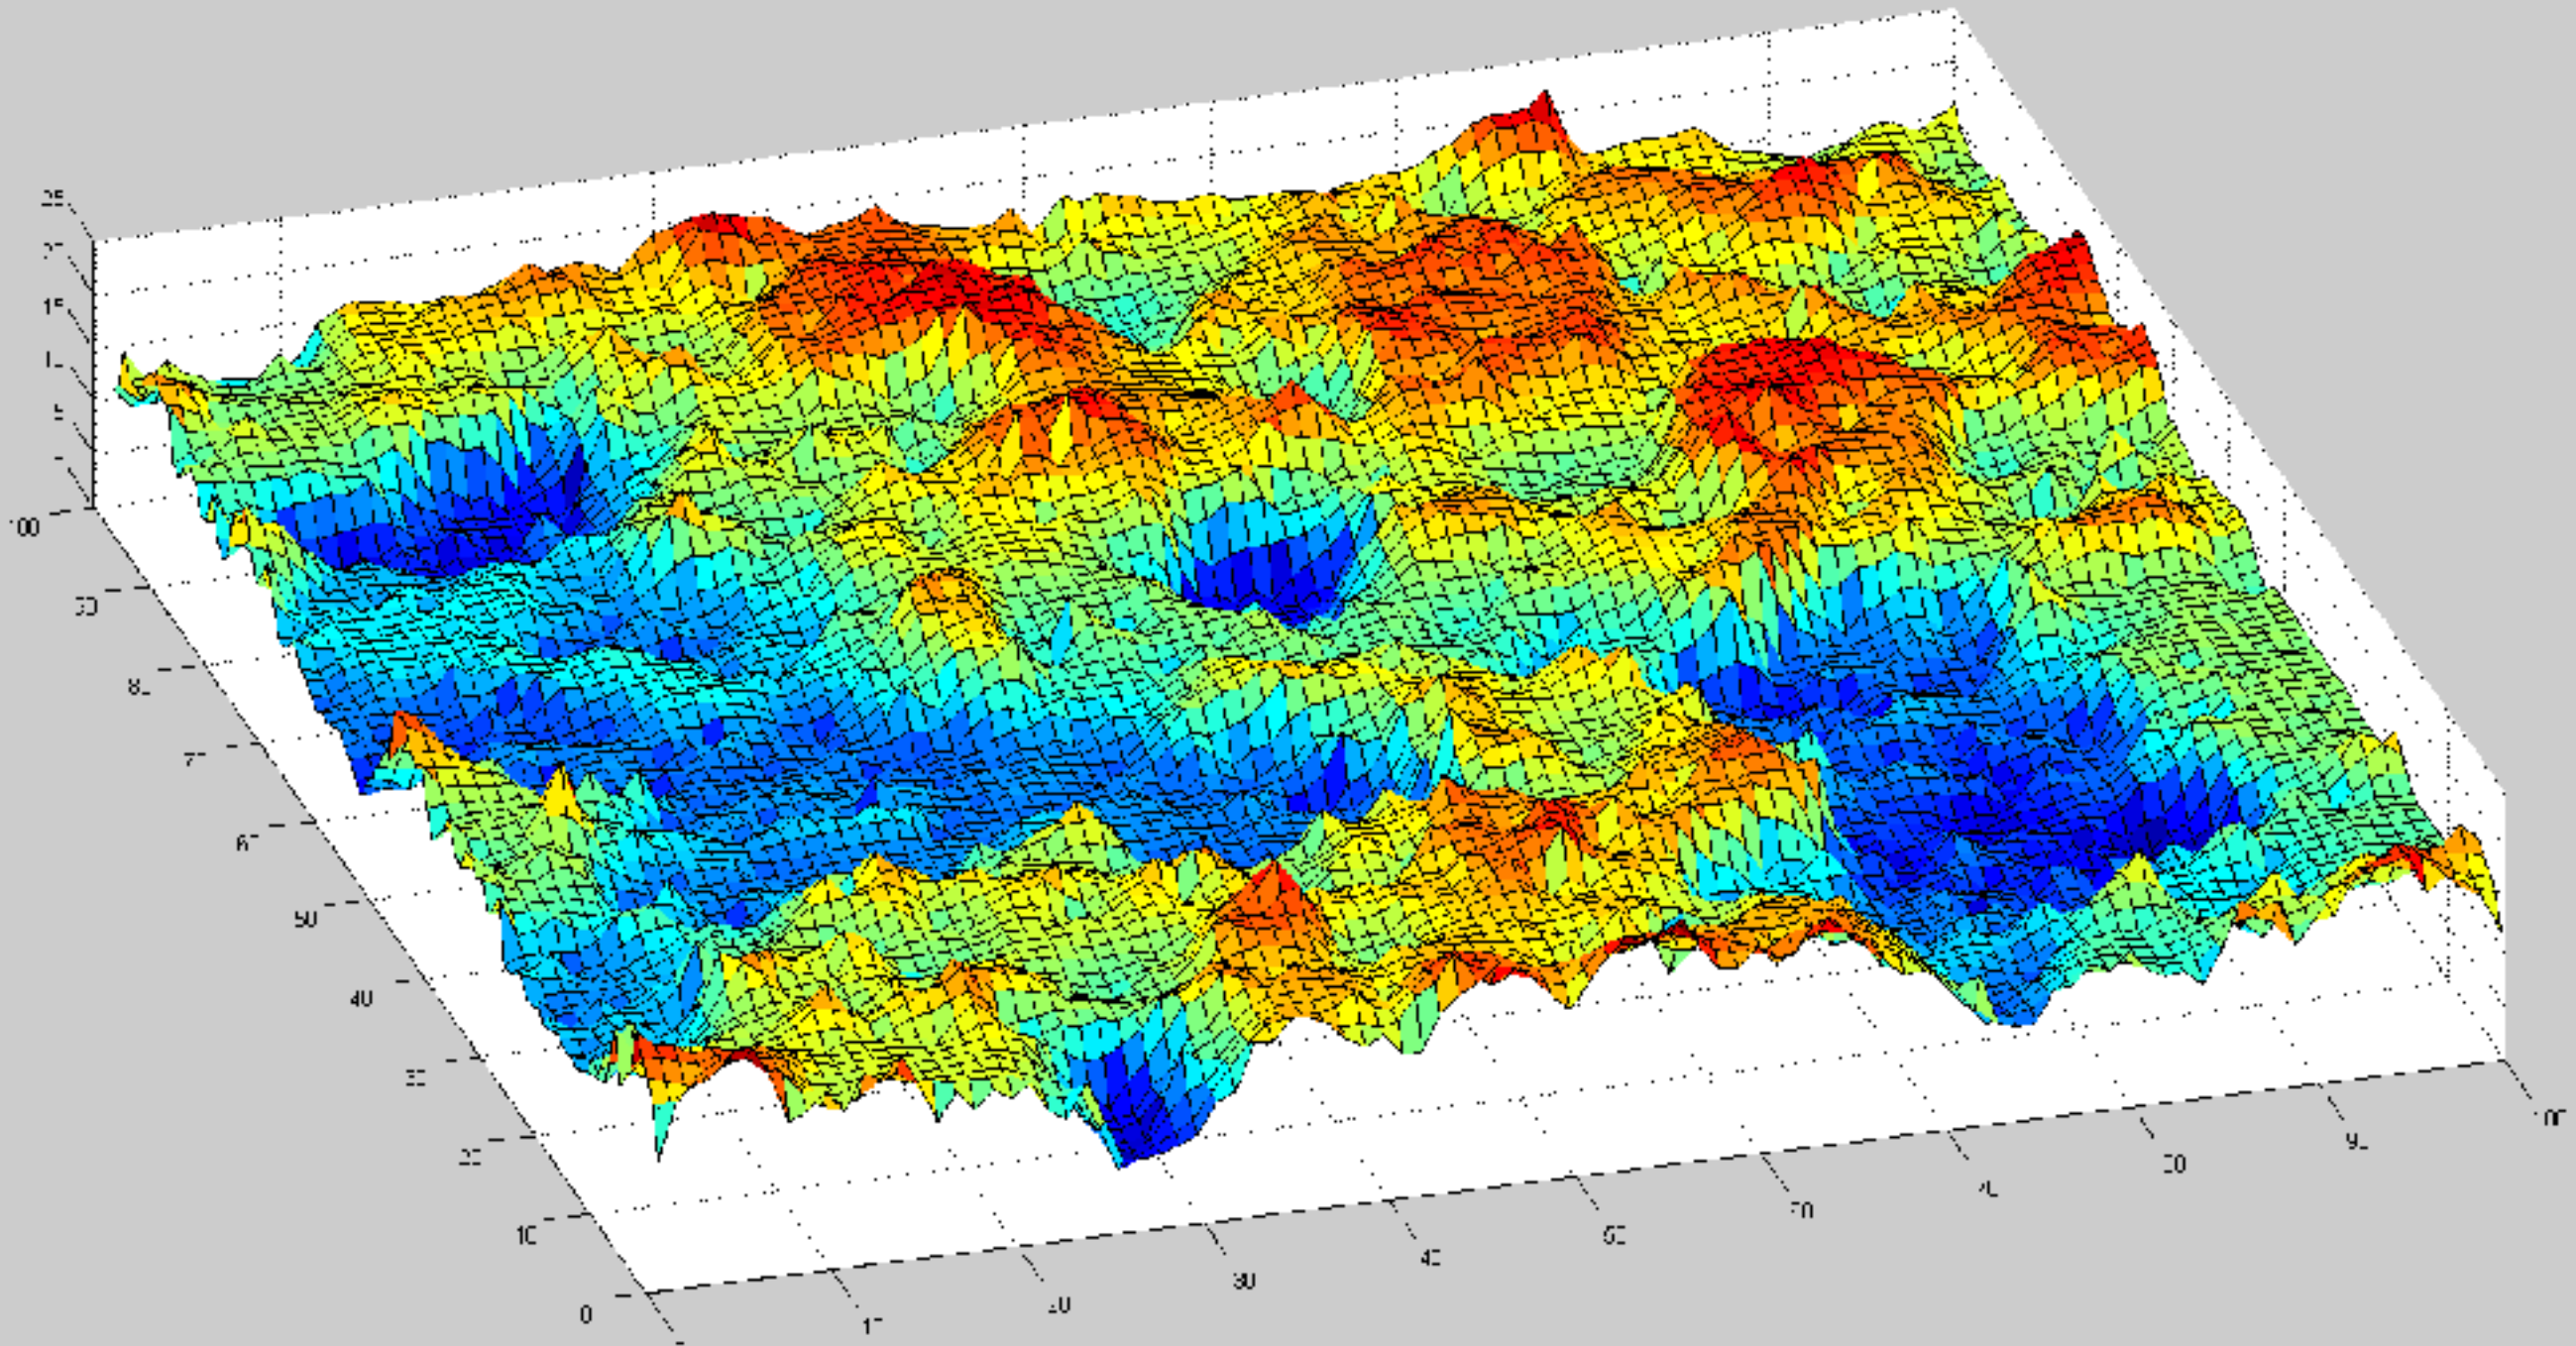
loop do

neighbor \leftarrow a highest-valued successor of *current*

if neighbor.VALUE \leq current.VALUE **then return** *current*.STATE

current \leftarrow *neighbor*


POTENTIAL PROBLEMS




► Retrieved from: <http://i.imgur.com/89KoR.png>

POTENTIAL PROBLEMS

- Local Maxima
- Plateau
- A flat plain



- Ridge
 - Any actions leads to lower ground, but combined actions lead to higher ground.



SOLVING PROBLEMS

- How to solve local minima, plateau, and ridge problems?

SOLVING PROBLEMS


- How to solve local minima, plateau, and ridge problems?
 - Random starting points
 - Backtracking
 - If we encounter the problem, just roll back some actions, and try a different one.
 - Simulated Annealing
 - Do not always choose the best actions. Probabilistically allow some bad actions and go with it. It may lead to a better solution.

GENETIC ALGORITHM

Local Search and Optimization Problem

INTRODUCING 8 QUEENS PROBLEM

.....



Retrieved from: <http://letstalkdata.com/2013/12/n-queens-part-1-steepest-hill-climbing/>

GENETIC ALGORITHM

Key Idea:

- Crossover and Mutation in DNA
- Natural Selection

Definition:

- Genetic Algorithm is a variant of stochastic beam search in which successor states are generated by combining two parent states rather than by modifying a single state.

GENETIC ALGORITHM: DEMO

Genetic Cars

http://rednuht.org/genetic_cars_2/

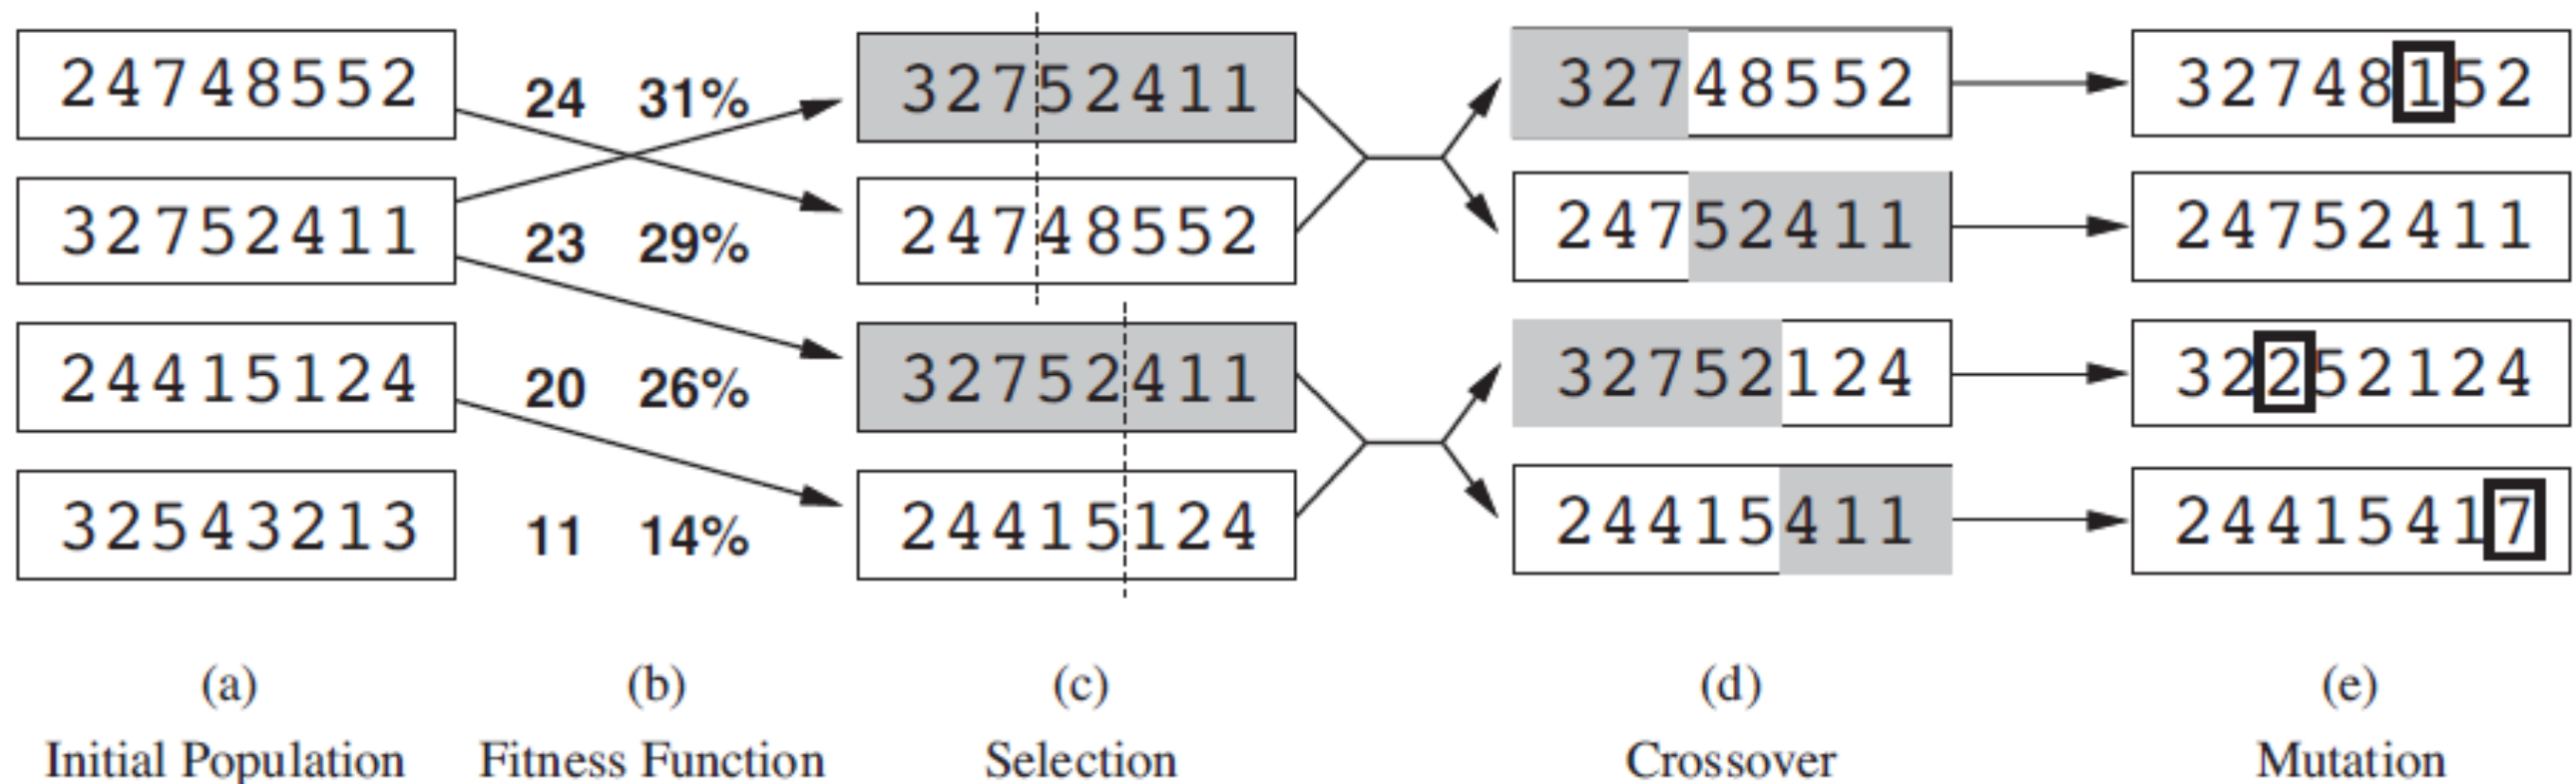
6 Species trying to reach the goal

[https://www.youtube.com/watch?](https://www.youtube.com/watch?v=2Q9bF8Ofyho&list=PLpJFWWWZXdN_qaI1SR6-7j6-5MyxhRqLQ0&index=2)

[v=2Q9bF8Ofyho&list=PLpJFWWWZXdN_qaI1SR6-7j6-5MyxhRqLQ0&index=2](https://www.youtube.com/watch?v=2Q9bF8Ofyho&list=PLpJFWWWZXdN_qaI1SR6-7j6-5MyxhRqLQ0&index=2)


GA: EXAMPLE

.....



GA: EXAMPLE

.....



function GENETIC-ALGORITHM(*population*, FITNESS-FN) **returns** an individual

inputs: *population*, a set of individuals

FITNESS-FN, a function that measures the fitness of an individual

repeat

new_population \leftarrow empty set

for $i = 1$ **to** SIZE(*population*) **do**

$x \leftarrow$ RANDOM-SELECTION(*population*, FITNESS-FN)

$y \leftarrow$ RANDOM-SELECTION(*population*, FITNESS-FN)

child \leftarrow REPRODUCE(x, y)

if (small random probability) **then** *child* \leftarrow MUTATE(*child*)

add *child* to *new_population*

population \leftarrow *new_population*

until some individual is fit enough, or enough time has elapsed

return the best individual in *population*, according to FITNESS-FN

function REPRODUCE(x, y) **returns** an individual

inputs: x, y , parent individuals

$n \leftarrow$ LENGTH(x); $c \leftarrow$ random number from 1 to n

return APPEND(SUBSTRING($x, 1, c$), SUBSTRING($y, c + 1, n$))

SEARCHING WITH NO OBSERVATION

WHAT IF...

The agent doesn't know where it is in the world.

- The agent knows only what it can do.
 - Know all action it can do.
 - Know all possible states.
- How to reach the goal state?

SEARCHING WITH NO OBSERVATION

Key idea:

- The robot doesn't know anything about positioning or environments (sensorless).
- It tries to reach the goal state.

Definition:

- When the agent's percepts provide no information at all, we have what is called a sensorless problem.

INTRODUCING VACUUM WORLD PROBLEM

1 Robot


2 Rooms

- Can be Clean/Dirty


3 Actions

- Left
- Right
- Suck


1




2




3




4




5




6



7



8









INTRODUCING VACUUM WORLD PROBLEM

Goal: clean both room.

Question:

- What actions does the robot have to take?

Group up, and try to derive the algorithm.








HOW TO SOLVE SENSORLESS PROBLEM

Find Belief States

- Every possible set of states, even if some states are unreachable from the initial state

Vacuum World

- There are 2 rooms. Each room is either clean or dirty
- Total belief states = $2*2*2 = 8$



HOW TO SOLVE SENSORLESS PROBLEM

We have all possible states.

Next step:


- Find possible action for each state.
 - Legal state: an agent can perform that action.
 - Illegal state: an agent cannot perform that action.
- Assumed that illegal state cannot break the world, we union all the possible actions.
- Assumed that illegal state can break the world, we intersect all the possible actions.

HOW TO SOLVE SENSORLESS PROBLEM


Original Action: Left,
Right, Suck

Assume that illegal
actions cannot break
the world, what are
the actions?


1




2




3



4




5




6



7



8









HOW TO SOLVE SENSORLESS PROBLEM

Original Action: Left,
Right, Suck

Assume that illegal
actions cannot break
the world, what are
the actions?

- Left
- Right
- Suck




HOW TO SOLVE SENSORLESS PROBLEM

We have states, and actions.

Next step:


- Transition Model
- Perform actions on states



HOW TO SOLVE SENSORLESS PROBLEM

Finally:

- Goal Test
 - Test if that state is the goal state or not
- Path Cost
 - Calculate cost of the path taken.



SEARCHING WITH PARTIAL OBSERVATION

SEARCHING WITH PARTIAL OBSERVATION


Key Idea:

- The agent observes the current state.
- The agent updates the knowledge base.
- The agent performs actions.
- The agent checks for the goal state.
- Loop back to the first state, if goal state is not reached.

INTRODUCING ROBOT LOCALIZATION PROBLEM

We don't know where the robot is


BUT the robot can report its environment !!



ROBOT LOCALIZATION PROBLEM

The robot just sent a signal NSW.


This means that the robot cannot move North, South, and West.



ROBOT LOCALIZATION PROBLEM

The robot just sent a signal NSW.


Let's update the map.



ROBOT LOCALIZATION PROBLEM

From the signal NSW, there are 4 possible locations.


Let's move the robot right.



ROBOT LOCALIZATION PROBLEM

The robot just sent another signal NS.


It means that the robot cannot move North and South.



ROBOT LOCALIZATION PROBLEM

The robot just sent a signal NS.


Let's update the map again.



ROBOT LOCALIZATION PROBLEM

We can also call the grid below an occupancy grid.

A grid of binary random variable in each cell. Each represents a presence of an obstacle.



ROBOT LOCALIZATION PROBLEM

Extending the algorithm we learnt so far, we can perform a Particle Filter.

We won't go into the mathematical details of Particle Filter.

However let's learn the concept:

<https://www.youtube.com/watch?v=aUkBa1zMKv4&t=247s>

MINIMAX

Game Search

MINIMAX

- Let's look at the animation of minimax first.
- <https://www.youtube.com/watch?v=KU9Ch59-4vw&t=63s>

MINIMAX

Key Idea:

- Playing games between 2 players: Checkers, Chess, ...
- Maximize outcomes of our plays, and minimize outcomes of opponent's plays


Definition:

- $\text{Minimax}(s) = \text{Utility}(s)$ if $\text{Terminal-Test}(s)$
 $= \max(\text{Minimax}(\text{Result}(s,a)))$ if $\text{Player}(s) = \text{Max}$
 $= \min(\text{Minimax}(\text{Result}(s,a)))$ if $\text{Player}(s) = \text{Min}$

MINIMAX: EXAMPLE

MAX


MIN



MINIMAX: EXAMPLE

MAX


MIN



MINIMAX: EXAMPLE

MAX


MIN



MINIMAX: EXAMPLE

MAX


MIN



MINIMAX: EXAMPLE

MAX


MIN



MINIMAX: EXAMPLE

MAX


MIN



MINIMAX: EXAMPLE

MAX


MIN



MINIMAX: EXAMPLE

MAX


MIN



MINIMAX: EXAMPLE

MAX

MIN



MINIMAX: THE ALGORITHM

.....

function MINIMAX-DECISION(*state*) **returns** *an action*
 return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$

function MAX-VALUE(*state*) **returns** *a utility value*
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$
 return *v*

function MIN-VALUE(*state*) **returns** *a utility value*
 if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow \infty$
 for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$
 return *v*

MULTIPLAYER GAME

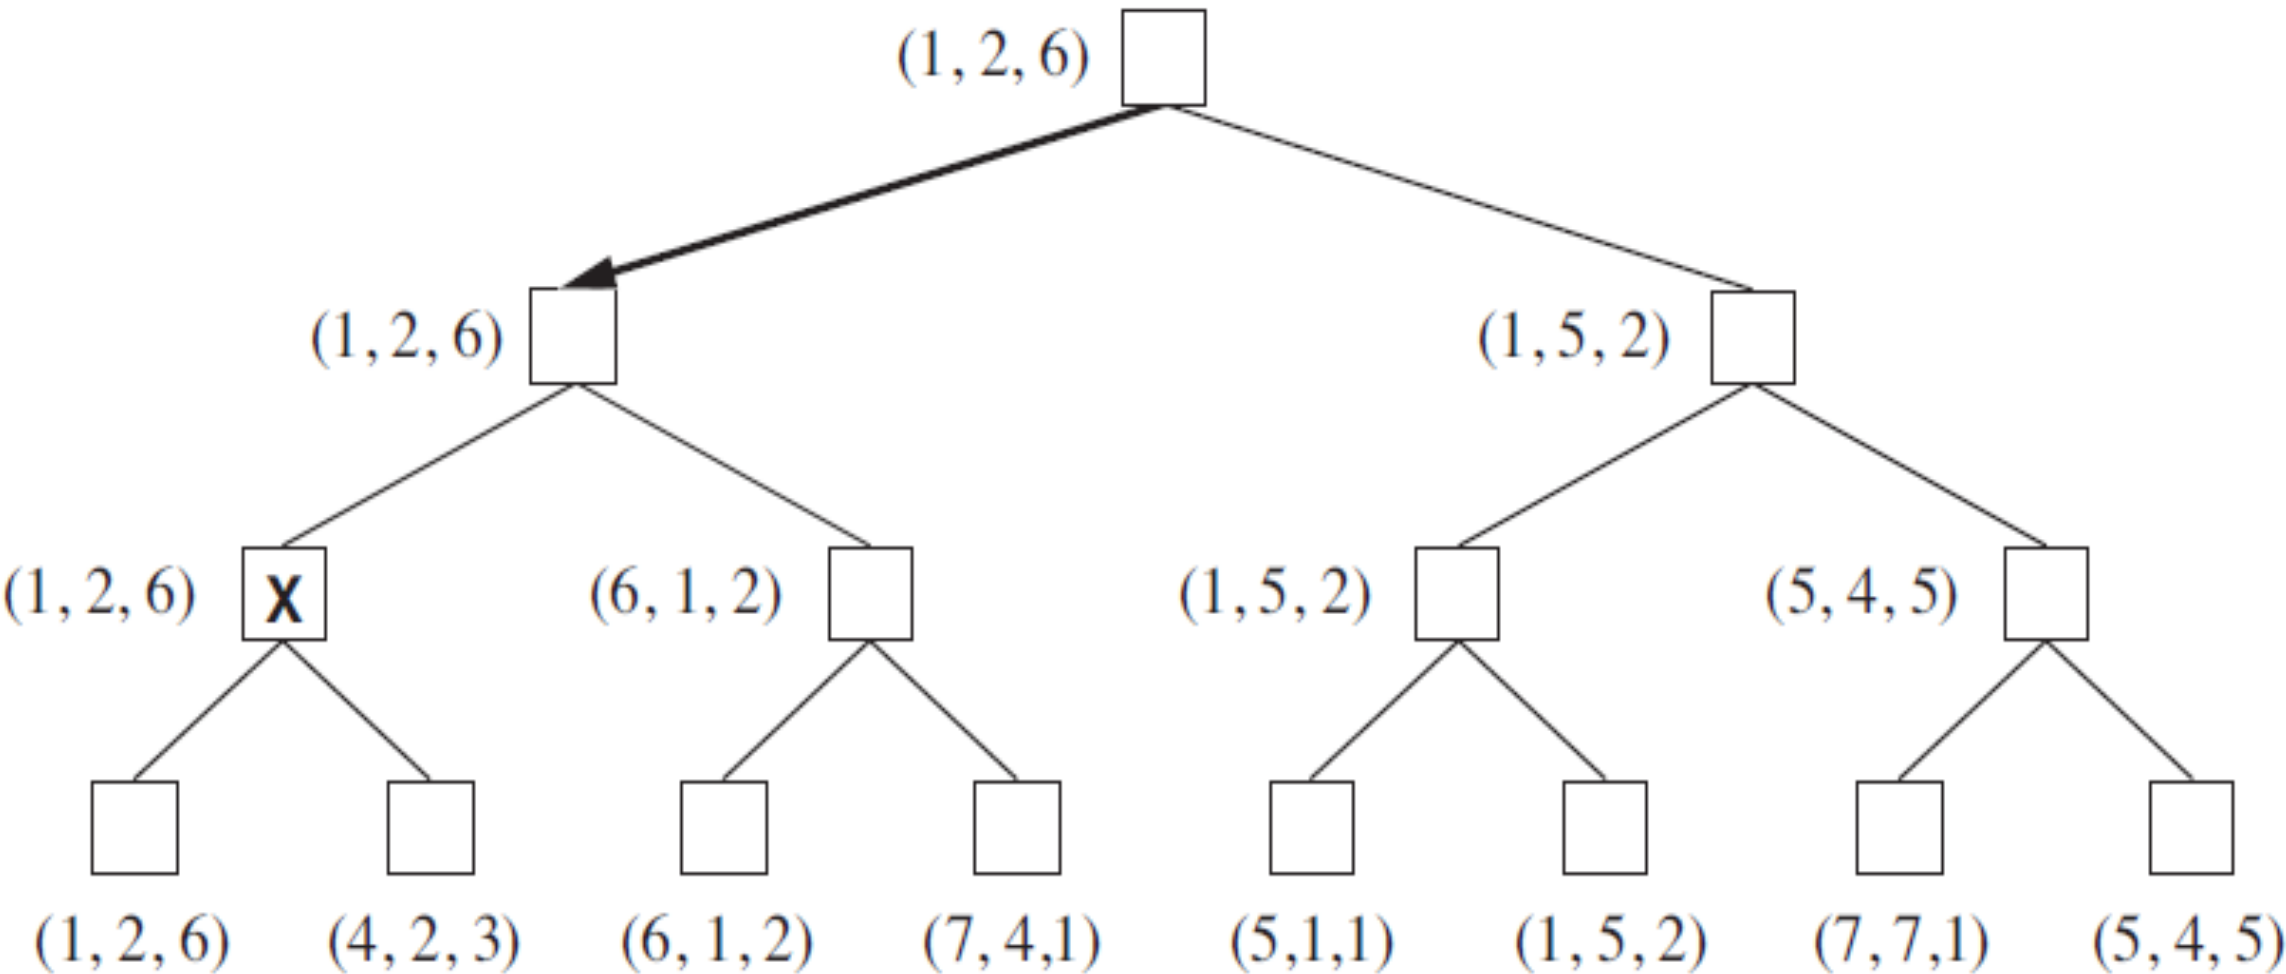
.....

to move
A

B

C

A



ONLINE SEARCH

Experience the Unknown

OFFLINE VS ONLINE

In offline search, we can compute a complete solution before entering the real world.

Online search computes actions on the go


- Take action
- Observe environment
- Goal test
- Plan next action

ONLINE SEARCH EXAMPLE

S = Starting Position

G = Goal State

We need to find an algorithm to “search to G.



ONLINE SEARCH: DFS ALGORITHM

function ONLINE-DFS-AGENT(s') **returns** an action
 inputs: s' , a percept that identifies the current state
 persistent: *result*, a table indexed by state and action, initially empty
 untried, a table that lists, for each state, the actions not yet tried
 unbacktracked, a table that lists, for each state, the backtracks not yet tried
 s, *a*, the previous state and action, initially null

if GOAL-TEST(s') **then return** *stop*
 if s' is a new state (not in *untried*) **then** *untried*[s'] \leftarrow ACTIONS(s')
 if s is not null **then**
 result[s , a] $\leftarrow s'$
 add s to the front of *unbacktracked*[s']
 if *untried*[s'] is empty **then**
 if *unbacktracked*[s'] is empty **then return** *stop*
 else $a \leftarrow$ an action b such that *result*[s' , b] = POP(*unbacktracked*[s'])
 else $a \leftarrow$ POP(*untried*[s'])
 $s \leftarrow s'$
 return a

ONLINE SEARCH EXAMPLE



function ONLINE-DFS-AGENT(s') **returns** an action

inputs: s' , a percept that identifies the current state

persistent: *result*, a table indexed by state and action, initially empty

untried, a table that lists, for each state, the actions not yet tried

unbacktracked, a table that lists, for each state, the backtracks not yet tried

s , a , the previous state and action, initially null

if GOAL-TEST(s') **then return** *stop*

if s' is a new state (not in *untried*) **then** $untried[s'] \leftarrow \text{ACTIONS}(s')$

if s is not null **then**

$result[s, a] \leftarrow s'$

add s to the front of $unbacktracked[s']$

if $untried[s']$ is empty **then**

if $unbacktracked[s']$ is empty **then return** *stop*

else $a \leftarrow$ an action b such that $result[s', b] = \text{POP}(unbacktracked[s'])$

else $a \leftarrow \text{POP}(untried[s'])$

$s \leftarrow s'$

return a

